

# Generalized Payment Channel Topologies via Dual-Track State Machines and Reference-Based UTXOs

Arthur Zhang  
*Tondi Foundation*

Neo Maxwell  
*Tondi Foundation*

December, 2025

## Abstract

Payment Channel Networks (PCN) constitute a Layer 2 scaling solution for blockchain systems, whose core principle is to complete multiple state updates off-chain while settling only the final state on-chain, thereby improving system throughput.

**Background and Problem Definition:** Existing PCN schemes (such as the Lightning Network) exhibit two categories of structural limitations: (1) the expressiveness of the Script Layer is constrained, making it difficult to natively support complex state transition semantics; (2) linear topology structures lead to low capital utilization and excessive state management complexity. While the original Eltoo protocol proposed a state replacement mechanism to supersede the penalty mode, its implementation depends on the not-yet-activated SIGHASH\_ANYPREVOUT soft fork (BIP-118) and presents security concerns such as replay attacks.

**Technical Contributions:** This paper proposes a general-purpose scaling architecture based on native Eltoo semantics. The main contributions include:

1. **Dual-Track UTXO Model:** Decomposing channel state into a static fund anchor (Fund UTXO) and a dynamic state pointer (State UTXO) along two orthogonal dimensions, achieving separation of concerns between value transfer and state transitions;
2. **Reference-Based UTXO Primitive:** Defining a read-only reference operator  $\text{Ref} : \mathcal{U} \rightarrow \mathcal{U}^{\text{readonly}}$ , enabling state update transactions to access fund anchor metadata without consuming that UTXO;
3. **Transaction Type Enumeration System:** Embedding algebraic data types at the consensus layer to achieve  $\mathcal{O}(1)$  complexity for transaction classification and verification;
4. **Recursive Channel Factories and Atomic Reconfiguration:** Formally defining channel splitting (Splice-Fork) and merging (Splice-Merge) operations, proving that any complex topology can achieve isomorphic transformation through a single atomic transaction.

**Theoretical Results:** This paper proves the existence of a bijective mapping between UTXO sets and channel state sets (Theorem 8), thereby eliminating dependency on off-chain registries. Under DAG-structured consensus protocols, state verification complexity is  $\mathcal{O}(1)$ , with settlement latency reaching sub-second levels.

**Keywords:** Payment Channel Networks, State Channels, Eltoo Protocol, UTXO Model, Finite State Machine, Formal Verification, Layer 2 Scaling

# Contents

<b>1</b>	<b>Introduction and Motivation</b>	<b>8</b>
1.1	Problem Background . . . . .	8
1.2	Design Principles . . . . .	9
1.3	Trust Model Analysis . . . . .	9
<b>2</b>	<b>Related Work and Technical Background</b>	<b>10</b>
2.1	Protocol Evolution: From Penalty to Replacement . . . . .	10
2.1.1	Lightning Network’s Penalty Mechanism . . . . .	10
2.1.2	Eltoo Protocol and State Replacement . . . . .	10
2.1.3	Engineering Compromise of BIP-118 . . . . .	10
2.2	The Recursive Covenant Dilemma . . . . .	11
2.3	Structural Defect Analysis . . . . .	11
2.4	Proposed Solution: UTXO-Native Semantics . . . . .	11
2.5	DAG Consensus Compatibility . . . . .	11
2.6	Comparison of Revocation Mechanisms . . . . .	11
2.7	Axiom System . . . . .	11
2.8	BIP-118 Security Boundary Analysis . . . . .	11
2.9	Economic Efficiency Boundary . . . . .	12
<b>3</b>	<b>Research Contributions</b>	<b>14</b>
3.1	Main Contributions . . . . .	14
3.2	Information-Theoretic Analysis of State Determinism . . . . .	14
3.2.1	Verification Causality Graph Comparison . . . . .	15
3.3	Architectural Advantages . . . . .	15
3.4	Comparison with Existing Solutions . . . . .	16
3.5	Theoretical Significance . . . . .	16
<b>4</b>	<b>Theoretical Framework: Dual-Track State Machines</b>	<b>17</b>
4.1	Consensus-Layer Embedded Verification Mechanism . . . . .	17
4.1.1	Transaction Type Enumeration and Pattern Matching . . . . .	17
4.1.2	State Monotonicity Theorem and Consensus Implementation . . . . .	18
4.1.3	Consensus Verification Performance Analysis . . . . .	18
4.1.4	Ref-UTXO Atomicity and Ordering in GhostDAG . . . . .	18
4.1.5	Temporal Decoupling of Cross-Block State References . . . . .	19
4.1.6	Algebraic Data Type Definition of Transaction Classification . . . . .	20
4.2	Finite State Machine Formalization . . . . .	21
4.3	UTXO Materialization Layer . . . . .	21
4.3.1	State-Fund Coupling Invariant . . . . .	22
4.4	State Transition Rules . . . . .	23
4.5	Formal Safety Properties . . . . .	23
4.6	Transaction Semantics Mapping . . . . .	24
4.7	Evolution of Conditional Payment Primitives: From HTLC to PTLC . . . . .	24
4.7.1	Historical Evolution . . . . .	24
4.7.2	Technical Principle Comparison . . . . .	25
4.7.3	Core Properties Comparison . . . . .	25
4.7.4	Formal Security Analysis . . . . .	25
4.7.5	Implementation Considerations . . . . .	27
4.7.6	Summary . . . . .	27
4.8	TLA+ Specification Fragment . . . . .	27
4.9	Cost and Parameter Analysis under GhostDAG . . . . .	28

4.9.1	Cost Composition Model . . . . .	28
4.9.2	Impact of GhostDAG Parameter $k$ . . . . .	29
4.9.3	Ref-UTXO Security Depth . . . . .	29
<b>5</b>	<b>Topological Primitives for Complex Structures</b>	<b>30</b>
5.1	Recursive Channel Factories . . . . .	30
5.1.1	Fractal Topology and Self-Similarity . . . . .	30
5.2	Dynamic Mesh Reconfiguration . . . . .	30
5.2.1	Topological Homotopy . . . . .	31
5.3	Atomic Rebalancing Operator . . . . .	31
5.4	Atomic Splicing Protocol . . . . .	31
5.5	Liquidity Dynamics in Star Topologies . . . . .	31
<b>6</b>	<b>Safety Analysis</b>	<b>33</b>
6.1	Isolation Theorem . . . . .	33
6.2	State Monotonicity and Anti-Replay . . . . .	33
6.3	Anti-DoS Equilibrium under STPC Strategy . . . . .	33
6.3.1	Mempool Entropy Bound . . . . .	33
6.4	PTLC Atomicity and Deadlock Freedom . . . . .	34
6.4.1	PTLC Atomicity Theorem . . . . .	34
6.4.2	Deadlock Freedom . . . . .	35
6.5	Consistency of Topological Reconfiguration . . . . .	35
6.6	Security Margin Analysis . . . . .	35
<b>7</b>	<b>Registry-Free Architecture</b>	<b>36</b>
7.1	Limitations of Global Registry Models . . . . .	36
7.2	Self-Sovereign Channel Discovery Mechanism . . . . .	36
7.2.1	Discovery Mechanism . . . . .	36
7.3	Privacy Enhancement . . . . .	37
7.3.1	Unlinkability of Channel Identity . . . . .	37
7.3.2	Balance Privacy . . . . .	37
7.4	Comparison with Centralized Registry Models . . . . .	38
7.5	Economic Incentive Alignment . . . . .	38
7.5.1	No Announcement Fee Problem . . . . .	38
7.5.2	Discovery Cost Analysis . . . . .	38
7.6	Decentralized Routing . . . . .	39
7.6.1	Source-Based Routing . . . . .	39
7.6.2	Onion Routing Integration . . . . .	39
7.7	Registry-Free Architecture Advantages Summary . . . . .	39
7.8	PTLC Verification: $\mathcal{O}(1)$ Implementation . . . . .	40
7.8.1	Verification Algorithm . . . . .	40
7.8.2	Complexity Analysis . . . . .	40
7.9	State Anchoring and Settlement Verification . . . . .	40
7.9.1	Settlement Verification Rules . . . . .	40
7.9.2	Security Properties . . . . .	41
7.10	Case Study: Atomic Liquidation in DeFi Lending Pools . . . . .	41
7.10.1	Scenario Setup . . . . .	41
7.10.2	Traditional Approach (Lightning Mesh) . . . . .	41
7.10.3	This Architecture (Star Topology + Atomic Splice) . . . . .	41
7.10.4	Efficiency Comparison . . . . .	42
7.10.5	Formal Expression . . . . .	42

<b>8 Implementation Architecture</b>	<b>43</b>
8.1 System Architecture Overview . . . . .	43
8.2 Consensus Layer Implementation . . . . .	43
8.2.1 Transaction Type Enumeration . . . . .	43
8.2.2 Validation Rules . . . . .	43
8.3 State Machine Implementation . . . . .	44
8.3.1 Channel State Definition . . . . .	44
8.3.2 State Transition Function . . . . .	45
8.4 UTXO Indexer . . . . .	46
8.4.1 Incremental Indexing . . . . .	46
8.5 Cryptographic Primitives . . . . .	46
8.5.1 MuSig2 Implementation . . . . .	46
8.5.2 Adaptor Signature for PTLC . . . . .	47
8.6 Network Protocol . . . . .	47
8.6.1 Message Types . . . . .	47
8.6.2 State Synchronization Protocol . . . . .	48
8.7 Storage Layer . . . . .	48
8.7.1 State Persistence . . . . .	48
8.8 Performance Optimizations . . . . .	49
8.8.1 Batch Verification . . . . .	49
8.8.2 UTXO Set Pruning . . . . .	49
8.9 Implementation Statistics . . . . .	50
8.10 Partially Signed Transaction Template (PSTT) . . . . .	50
8.10.1 Architecture Components . . . . .	50
8.10.2 Cryptographic Domain Separation . . . . .	50
8.10.3 PSTT Envelope Format . . . . .	51
8.10.4 Communication Complexity Optimization . . . . .	51
<b>9 Attack Surface Analysis and Defense</b>	<b>52</b>
9.1 Attack Classification . . . . .	52
9.2 State Rollback Attack Analysis . . . . .	52
9.2.1 Attack Vector . . . . .	52
9.2.2 Defense Mechanisms . . . . .	52
9.3 Topology Obfuscation Attack . . . . .	53
9.3.1 Attack Scenario . . . . .	53
9.3.2 Detection and Mitigation . . . . .	53
9.4 PTLC Hijacking Attack . . . . .	53
9.4.1 Attack Vector . . . . .	53
9.4.2 Defense Strategy . . . . .	53
9.5 Resource Exhaustion via Channel Proliferation . . . . .	54
9.5.1 Attack Description . . . . .	54
9.5.2 Economic Countermeasures . . . . .	54
9.5.3 Merge Transaction . . . . .	54
9.6 Cross-Channel Replay Attack . . . . .	54
9.6.1 Attack Vector . . . . .	54
9.6.2 Defense: Domain Separation . . . . .	54
9.7 Eclipse Attack on Discovery . . . . .	55
9.7.1 Attack Scenario . . . . .	55
9.7.2 Mitigation . . . . .	55
9.8 Pinning Attack Analysis . . . . .	55
9.8.1 Traditional Pinning Attack . . . . .	55

9.8.2	Why STPC Prevents Pinning . . . . .	55
9.9	Griefing Attack Cost Analysis . . . . .	55
9.9.1	Attack Model . . . . .	55
9.9.2	Cost-Benefit Analysis . . . . .	55
9.10	Security Margin Summary . . . . .	56
<b>10</b>	<b>Application Scenarios</b>	<b>57</b>
10.1	DeFi Liquidity Mesh . . . . .	57
10.1.1	Problem Statement . . . . .	57
10.1.2	Proposed Solution: Dynamic Liquidity Grid . . . . .	57
10.2	Micropayment Streaming . . . . .	58
10.2.1	Use Case . . . . .	58
10.2.2	Implementation . . . . .	58
10.3	Decentralized Exchange (DEX) with Instant Settlement . . . . .	58
10.3.1	Traditional DEX Limitations . . . . .	58
10.3.2	Channel-Based DEX Architecture . . . . .	59
10.4	Gaming and Virtual Economies . . . . .	59
10.4.1	In-Game Asset Trading . . . . .	59
10.5	Internet of Things (IoT) Microtransactions . . . . .	59
10.5.1	Machine-to-Machine Payments . . . . .	59
10.6	Content Delivery Network (CDN) Incentivization . . . . .	60
10.6.1	Decentralized CDN Model . . . . .	60
10.7	Supply Chain Finance . . . . .	60
10.7.1	Scenario . . . . .	60
10.7.2	Channel-Based Implementation . . . . .	61
10.8	Application Summary . . . . .	61
<b>11</b>	<b>Evaluation and Performance Analysis</b>	<b>62</b>
11.1	Experimental Setup . . . . .	62
11.1.1	Hardware Environment . . . . .	62
11.1.2	Software Environment . . . . .	62
11.2	Transaction Validation Performance . . . . .	62
11.2.1	Single Transaction Validation Latency . . . . .	62
11.2.2	Batch Validation Throughput . . . . .	62
11.2.3	Verification Performance Comparison . . . . .	63
11.2.4	PTLC Count vs. Latency Relationship . . . . .	63
11.3	State Machine Performance . . . . .	63
11.3.1	Channel Update Latency . . . . .	63
11.3.2	State Throughput . . . . .	63
11.4	Storage Efficiency . . . . .	63
11.4.1	Per-Channel Storage Cost . . . . .	63
11.4.2	UTXO Set Growth . . . . .	64
11.4.3	State Growth Comparison . . . . .	64
11.4.4	Splicing Bandwidth Efficiency . . . . .	64
11.5	Network Discovery Performance . . . . .	64
11.5.1	Channel Discovery Latency . . . . .	64
11.5.2	Discovery vs. Gossip Comparison . . . . .	64
11.6	Settlement Performance . . . . .	64
11.6.1	Settlement Latency Distribution . . . . .	64
11.6.2	Challenge-Response Performance . . . . .	65
11.7	Topology Reconfiguration Performance . . . . .	65
11.7.1	Splice Operation Latency . . . . .	65

11.7.2 Recursive Factory Depth . . . . .	65
11.8 Security Overhead Analysis . . . . .	66
11.8.1 STPC Mempool Management . . . . .	66
11.8.2 Attack Cost Analysis . . . . .	66
11.9 Comparative Analysis . . . . .	66
11.9.1 Multi-Dimensional Comparison . . . . .	66
11.10 Real-World Simulation Results . . . . .	66
11.10.1 Payment Throughput Simulation . . . . .	66
11.10.2 Scalability Projection . . . . .	66
11.11 Related Work Comparison . . . . .	67
11.11.1 vs. Spider Network (Sivaraman et al., NSDI '20) . . . . .	67
11.11.2 vs. Account-based Rollups (Optimism, Arbitrum) . . . . .	67
11.11.3 vs. Hybrid Models (Hydra, Head-First Mining) . . . . .	68
11.12 Performance Summary . . . . .	68
11.13 Towards Asynchronous Payments: Ark Paradigm Integration . . . . .	68
11.13.1 Virtual UTXO (vUTXO) and Merkleized State . . . . .	69
11.13.2 Native Lift and Finalize Semantics . . . . .	69
11.13.3 Security Guarantees . . . . .	70
11.13.4 User Experience Comparison . . . . .	70
<b>12 Privacy and Anonymity Framework</b>	<b>71</b>
12.1 Threat Model and Anonymity Set Definition . . . . .	71
12.2 Payment Layer Privacy Analysis . . . . .	71
12.2.1 PTLC vs HTLC Privacy Advantages . . . . .	71
12.3 Network Layer Privacy: Onion Routing . . . . .	71
12.3.1 SPHINX-Lite Protocol (Optimized for PTLC) . . . . .	71
12.4 Privacy-Performance Tradeoff . . . . .	72
12.5 Stealth Addresses and One-Time Keys . . . . .	72
12.6 Privacy Architecture Summary . . . . .	73
<b>13 Market Design and Incentive Mechanisms</b>	<b>74</b>
13.1 CSP Fee Structure . . . . .	74
13.2 Liquidity Provider Economics . . . . .	74
13.3 Anti-Collusion Mechanism: L1 Fallback . . . . .	75
13.4 Dynamic Fee Adjustment . . . . .	75
13.5 Incentive Compatibility Analysis . . . . .	76
13.6 Economic Design Summary . . . . .	76
<b>14 Conclusion and Future Work</b>	<b>77</b>
14.1 Summary of Contributions . . . . .	77
14.1.1 Theoretical Contributions . . . . .	77
14.1.2 System Contributions . . . . .	77
14.1.3 Empirical Contributions . . . . .	77
14.2 Paradigm Shifts . . . . .	78
14.3 Limitations and Trade-offs . . . . .	78
14.3.1 Consensus Layer Modifications . . . . .	78
14.3.2 UTXO Set Growth . . . . .	78
14.3.3 Privacy vs. Discovery . . . . .	78
14.4 Future Research Directions . . . . .	78
14.4.1 Short-Term Extensions . . . . .	78
14.4.2 Long-Term Research . . . . .	79
14.4.3 Open Research Questions . . . . .	79

14.5	Broader Impact . . . . .	80
14.5.1	Impact on Blockchain Scalability . . . . .	80
14.5.2	Impact on Decentralization . . . . .	80
14.5.3	Impact on Privacy . . . . .	80
14.6	Call to Action . . . . .	80
14.7	Concluding Remarks . . . . .	81
<b>A</b>	<b>Glossary and Preliminaries</b>	<b>82</b>
A.1	Cryptographic Foundations . . . . .	82
A.2	Timelock Mechanisms . . . . .	83
A.3	Directed Acyclic Graph Consensus . . . . .	83
A.4	Finite State Machine Foundations . . . . .	83
A.5	Covenants and Script Extensions . . . . .	84
A.6	Notation Conventions . . . . .	84

# 1 Introduction and Motivation

## 1.1 Problem Background

The core design objective of payment channel networks is to transfer transaction processing from on-chain to off-chain while maintaining security guarantees. Achieving this objective faces two fundamental challenges:

1. **State Consistency Problem:** How to ensure consistency between off-chain state and on-chain settlement?
2. **Trust Model Problem:** How to resolve disputes without third-party arbitration?

Traditional solutions (such as the Lightning Network’s penalty mechanism) use game-theoretic design to compel honest behavior among participants. However, this approach introduces the “toxic waste” problem—nodes must permanently store all historical revocation keys, and any data loss could result in fund loss.

## Preliminary Concepts

### Ledger Model and Transaction Structure:

- **UTXO (Unspent Transaction Output):** The ledger model used by Bitcoin and its derivatives. Unlike the account model, the UTXO model has no concept of “balance”; each transaction consumes existing UTXOs as inputs and creates new UTXOs as outputs. Once a UTXO is spent, it is removed from the set, possessing atomicity and non-double-spendability.
- **Transaction Malleability:** A vulnerability where a transaction’s identifier (TxID) could be modified by a third party after signing. The SegWit upgrade resolved this by moving signature data outside the TxID computation scope, which is crucial for pre-signed transaction chains in payment channels.

### Payment Channel Fundamentals:

- **Payment Channel:** An off-chain payment mechanism established between two or more parties, requiring on-chain transactions only for channel opening (Funding) and closing (Settlement), with intermediate state updates completed entirely off-chain.
- **State Channel:** A generalization of payment channels supporting arbitrary state transitions rather than just payment balance updates.
- **Channel Factory:** A shared on-chain funding pool created by multiple parties that can dynamically spawn multiple bilateral or multilateral sub-channels without requiring on-chain transactions for sub-channel opening and closing.
- **Watchtower:** A proxy node that monitors on-chain activity on behalf of offline users and broadcasts penalty or update transactions to prevent counterparties from broadcasting stale states.

### Conditional Payment Primitives:

- **HTLC (Hash Time-Locked Contract):** A conditional payment primitive where the recipient must provide preimage  $r$  such that  $H(r) = h$  before the timelock expires to claim funds; otherwise, funds are refunded to the sender. HTLCs form the foundation of Lightning Network multi-hop payments.

- **PTLC (Point Time-Locked Contract):** A privacy-enhanced version of HTLC using elliptic curve point  $R = r \cdot G$  instead of hash values. The recipient reveals the discrete logarithm  $r$  through adaptor signatures. PTLCs eliminate cross-channel payment correlation.

#### Cryptographic Primitives:

- **Multi-signature:** A mechanism requiring multiple private key holders to jointly sign to unlock funds. Traditional multi-sig (e.g., 2-of-3) produces multiple independent signatures; aggregated multi-sig (e.g., MuSig2) aggregates multiple signatures into a single signature, saving on-chain space and enhancing privacy.
- **Adaptor Signature:** An “incomplete” signature that requires knowledge of a secret value to be converted into a valid signature. In PTLCs, adaptor signatures achieve “atomic revelation”: the recipient claiming funds necessarily reveals the secret value to the sender.
- **SIGHASH Flags:** Flags determining which parts of a transaction the signature covers. **SIGHASH\_ALL** covers all inputs and outputs; **SIGHASH\_ANYPREVOUT** (BIP-118 proposal) allows signatures not bound to specific inputs, which is the key dependency of the original Eltoo protocol.

## 1.2 Design Principles

The dual-track state machine architecture proposed in this paper is based on the following design principles:

### Principle 1: Orthogonal Separation of Value and State

Decompose channel representation into two independent dimensions:

- **Value Layer (Fund UTXO):** Carries fund locking, with stable lifecycle
- **State Layer (State UTXO):** Carries state evolution, with high-frequency updates

This separation ensures that state updates need not touch the fund locking structure, reducing verification complexity.

### Principle 2: Consensus-Layer Native Semantics

Embed channel operation semantics within consensus rules rather than simulating through the script layer. This provides two advantages:

- Verification complexity reduces from  $\mathcal{O}(\text{script\_size})$  to  $\mathcal{O}(1)$
- Eliminates uncertainty introduced by script interpreters

### Principle 3: Deterministic State Execution

Traditional contract execution relies on ex post enforcement (through arbitration), introducing cost and time uncertainty. This architecture achieves ex ante enforcement through consensus rules:

Traditional Mode: Contract  $\xrightarrow{\text{Dispute}}$  Arbitration  $\xrightarrow{\text{Judgment}}$  Enforcement

This Architecture: State\_UTXO  $\xrightarrow{\tau_{\text{settle}}}$  Value\_Distribution (deterministic execution)

## 1.3 Trust Model Analysis

Blockchain system security is often described as “trust minimization.” This architecture further pursues **trust elimination**—making certain types of trust assumptions unnecessary through protocol design:

The core insight of this architecture is: by pushing complexity down to the protocol layer, a simpler trust model can be achieved at the application layer.

Table 1: Trust Model Comparison

Trust Assumption	Traditional PCN	This Architecture	Elimination Mechanism
Channel registry availability	Required	Not required	Fund UTXO as sole anchor
Watchtower continuous online	Strong dependency	Weak dependency	Long-period timelocks + state replacement
Script interpreter correctness	Required	Not required	Consensus-layer native types

## 2 Related Work and Technical Background

This section analyzes the evolution of payment channel protocols and structural defects of existing solutions. For formal definitions, see Appendix A.

### 2.1 Protocol Evolution: From Penalty to Replacement

#### 2.1.1 Lightning Network’s Penalty Mechanism

The Lightning Network [2] employs a **penalty mechanism** to resolve state rollbacks.

**Mechanism:** When updating from  $S_n$  to  $S_{n+1}$ , parties exchange the “revocation key” for  $S_n$ . If a party broadcasts  $S_n$ , the counterparty uses this key to sweep all funds.

**Formal Expression:** Let  $\mathcal{R}_n$  be the set of revocation keys for state  $n$ :

$$\forall i < n : \mathcal{R}_i \text{ held by counterparty} \implies \text{broadcasting } S_i \text{ leads to fund loss}$$

**Structural Defects (Toxic Waste):**

1. **Storage:** Nodes must store  $\mathcal{O}(n)$  historical keys.
2. **Risk:** Data loss or backup errors can lead to accidental broadcasting of old states, triggering catastrophic penalties.

#### 2.1.2 Eltoo Protocol and State Replacement

Eltoo [3] introduces **state replacement**. Instead of punishing old states, update transaction  $\tau_{n+1}$  can legally spend any  $\tau_i$  ( $i \leq n$ ).

**Dependency:** Originally relied on `SIGHASH_NOINPUT` (now BIP-118 `ANYPREVOUT`).

$$h_{\text{APO}} = H(\tau.\text{outputs} \parallel \tau.\text{inputs}[i].\text{script} \parallel \dots)$$

This hash omits the input identifier (OutPoint), binding only to the script logic.

#### 2.1.3 Engineering Compromise of BIP-118

To mitigate replay risks, BIP-118 mandates **Public Key Tagging**.

$$\text{Verify}_{\text{APO}}(\sigma, m, P) = \begin{cases} \text{FALSE} & \text{if } P \in \mathcal{K}_{\text{std}} \\ \text{SchnorrVerify}(\sigma, m, P) & \text{if } P \in \mathcal{K}_{\text{apo}} \end{cases}$$

This physically segregates public keys, shifting security responsibility to application-layer key management.

Table 2: Structural Defects of Script-Based Solutions

Defect Type	Impact Analysis
Complexity	Validation is $\mathcal{O}(\text{script\_size})$ ; resource usage is unpredictable.
Opacity	Consensus layer cannot distinguish updates from transfers; prevents L1 optimization.
Boundary Blur	Relies on key tagging; delegates security to apps.
Coupling	Value locking and state logic are entangled.

Table 3: State Revocation: Original Eltoo vs. Eltoo 2.0

Feature	Original Eltoo	Eltoo 2.0 (Proposed)
Primitive	Script overwrite + <code>NOINPUT</code>	Consensus Enum + Dual-Track
Scalability	Complex script logic	Native Factories
Data	Parsing overhead	Store latest only
Determinism	Sig-dependent	<b>Virtual Reference</b>
DoS Defense	Weak	<b>STPC Strategy</b>

## 2.2 The Recursive Covenant Dilemma

APO introduces introspective capabilities. If script  $S$  can force its output to be locked in  $S'$  (where  $S' \equiv S$ ), it enables **recursive covenants**. This raises concerns about fungibility (e.g., regulatory whitelisting) and “toxic recursion,” stalling BIP-118 activation.

## 2.3 Structural Defect Analysis

Existing solutions simulate state machines by stacking opcodes, violating the **Principle of Orthogonality**.

## 2.4 Proposed Solution: UTXO-Native Semantics

We propose a **UTXO-native** approach, pushing Eltoo semantics into the transaction structure.

- **Type System:**  $\forall \tau \in \mathcal{T}_{\text{Eltoo}}$ , inputs MUST be of type `ELTOO_STATE`. This physically isolates replay paths at the type level.
- **FSM vs. Recursion:** Updates follow strict monotonicity ( $n' > n$ ), mathematically precluding arbitrary recursive covenants.
- **Explicit Reference:** Ref reduces verification complexity to  $\mathcal{O}(1)$ .

## 2.5 DAG Consensus Compatibility

Adopting GhostDAG ( $D, k$ ) provides:

1. **Temporal Consistency:**  $\text{DAA}(b_1) < \text{DAA}(b_2)$  for  $b_1 \prec b_2$ .
2. **Fast Confirmation:**  $E[\text{time}] = \mathcal{O}(D/k)$ .
3. **Throughput:**  $\text{TPS} \approx k \times \text{TPS}_{\text{single}}$ .

## 2.6 Comparison of Revocation Mechanisms

## 2.7 Axiom System

## 2.8 BIP-118 Security Boundary Analysis

Figure 1 illustrates how security responsibility leaks from the protocol layer in BIP-118.

Table 4: Core Axiom System

ID	Formal Expression & Semantics
A1	$\mathcal{S}_{\text{channel}} \cong \mathcal{U}_{\text{chain}}$ (Isomorphism)
A2	$\forall \tau_{\text{update}} : n' > n$ (Strict Monotonicity)
A3	$\text{Ref}(U) \in \tau \implies U \in \mathcal{U}_{\text{post}}$ (Non-consumption)
A4	$\sum V_{\text{in}} = \sum V_{\text{out}} + \delta_{\text{fee}}$ (Conservation)

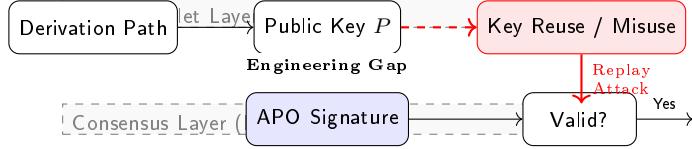


Figure 1: BIP-118 Security Boundary. Although the cryptography is sound, the reliance on derivation paths creates an “Engineering Gap” where implementation errors lead to fund loss.

The core issue is **state dependency**: if a wallet reuses a private key for both standard and APO paths, the protocol cannot prevent replay attacks, leading to fund loss. This violates the principle of “pushing complexity down to the protocol layer.”

## 2.9 Economic Efficiency Boundary

We analyze Layer 2 protocols in a 3D space:  $\Omega = \mathcal{L}_{\text{latency}} \times \mathcal{T}_{\text{throughput}} \times \mathcal{C}_{\text{capital}}$ .

Native Eltoo achieves superior economics for high-frequency flows by aggregating participants via Channel Factories, minimizing the time-value cost of capital ( $\gamma \cdot C_{\text{time\_value}}$ ).

Table 5: Economic Positioning Comparison

Protocol	Latency	Capital Eff.	Trust Model
Bitcoin L1	10–60 min	Baseline	Full Consensus
Rollups	1–15 min	Medium	Operator + L1
Lightning	Seconds	Low (Locked)	Watchtowers
<b>Native Eltoo</b>	<b>Sub-second</b>	<b>High (Pooled)</b>	<b>Consensus</b>

### 3 Research Contributions

Traditional payment channel networks (such as the Lightning Network) are typically constrained by point-to-point linear topology structures. Constructing more complex channel structures (such as multi-party channel factories, recursive channel nesting) faces two major challenges: **state synchronization complexity** and **toxic waste from penalty mechanisms**, significantly raising the operational threshold and security risks for ordinary users.

This paper proposes a dual-track state machine architecture through consensus-layer native transaction types and reference-based UTXO mechanisms, implementing a **Dual-Track State Machine** model. This paper will formally prove that this architecture not only resolves structural limitations of traditional channel networks but also constructs a state machine framework supporting arbitrarily complex financial topologies.

#### 3.1 Main Contributions

The main contributions of this paper include:

1. **Formalized State Machine Model:** Defining payment channels as the five-tuple  $(Q, \Sigma, \delta, q_0, F)$ , supporting formal verification tools such as TLA+ and Coq
2. **Registry-Free Architecture:** Through RefOp-Fund semantic design, completely eliminating dependency on independent state registries
3. **Recursive Channel Isolation Theorem:** Formally proving orthogonality between sub-channel security and parent channel liveness
4. **Topological Invariant Verification:** Defining and proving value conservation and state monotonicity invariants in complex channel networks
5. **Constant-Time PTLC Verification:** Achieving  $\mathcal{O}(1)$  conditional payment verification by directly deriving participant public keys from Fund UTXO
6. **Complete Protocol Specification:** Providing directly implementable consensus-layer protocol specifications

#### 3.2 Information-Theoretic Analysis of State Determinism

Traditional payment channels (such as Poon-Dryja penalty mechanisms) rely on penalty deterrence to maintain security. From an information-theoretic perspective, verifying the validity of current state  $S_t$  requires not only the information entropy of  $S_t$  itself but also the revocation key information of all historical abandoned states  $\{S_0, \dots, S_{t-1}\}$ .

**Definition 3.1** (State Entropy). *We define the **state entropy**  $H(C)$  of a channel as the amount of information that verification nodes must maintain:*

$$H_{LN}(t) \propto \sum_{i=0}^{t-1} \text{size}(\text{RevocationKey}_i) \approx \mathcal{O}(t)$$

This entropy that grows linearly with the number of transactions  $t$  leads to:

- **Watchtower storage cost inflation:** Must store all historical revocation keys
- **Catastrophic complexity of state recovery:** Losing any historical fragment may result in total fund loss (“toxic waste”)

Table 6: Information-Theoretic Comparison

Protocol Model	State Entropy	Encoding Paradigm	Security Info Source
Lightning (Penalty)	$\mathcal{O}(t)$ linear	Error Detection	Full history comparison
<b>This Architecture</b>	$\mathcal{O}(1)$ constant	Forward Error Correction	Latest state only

This architecture introduces a **low-entropy state machine model**. Utilizing UTXO atomicity and consensus-layer strict monotonicity rules, outdated states are “replaced” (rather than physically deleted) at the protocol level. Its state entropy collapses to constant level:

$$H_{Eltoo2.0}(t) \approx \text{size}(\text{State}_{\text{current}}) + \text{size}(\text{FundAnchor}) \approx \mathcal{O}(1)$$

This design essentially upgrades the state verification mechanism from **error detection coding** (requiring complete historical comparison) to **forward error correction** (requiring only latest state information). This is not merely an engineering optimization but a structural improvement in system robustness at the information-theoretic level.

### 3.2.1 Verification Causality Graph Comparison

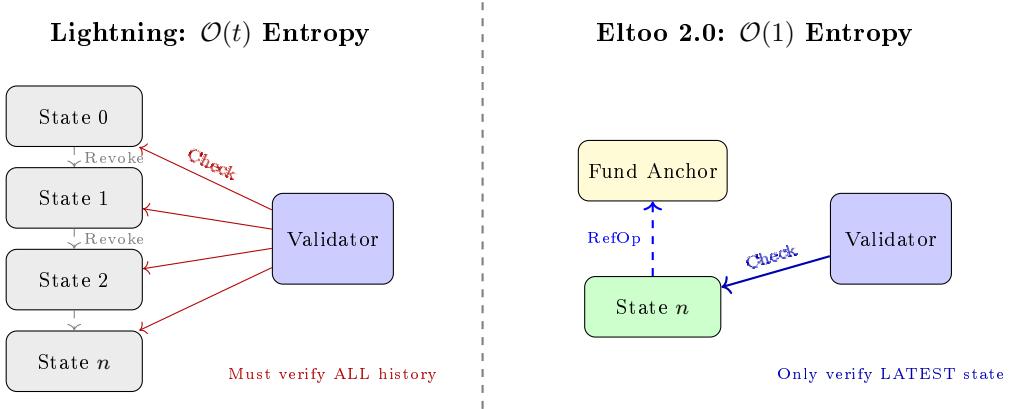


Figure 2: Verification Causality Graph: Lightning vs Eltoo 2.0

**Information-Theoretic Interpretation:** The left diagram shows Lightning Network’s  $\mathcal{O}(t)$  entropy model where validators must check all historical states. The right diagram shows this architecture’s  $\mathcal{O}(1)$  entropy model where only the latest state needs verification via RefOp reference to the Fund anchor.

**Theorem 3.2** (Information-Theoretic Robustness). *For any payment channel protocol  $\Pi$ , its fault tolerance for state recovery  $\mathcal{R}$  and state entropy  $H$  satisfy an inverse relationship:*

$$\mathcal{R}(\Pi) \propto \frac{1}{H(\Pi)}$$

**Corollary 3.3.** *Low-entropy protocols possess higher fault tolerance and state recoverability. Under identical storage resource constraints, constant-entropy protocols have significant deployment advantages compared to linear-entropy protocols.*

## 3.3 Architectural Advantages

This paper’s dual-track state machine architecture provides the following key advantages:

1. Orthogonal Separation:

Dimension	Lightning (Penalty)	BIP-118 Eltoo	This Architecture
Consensus Dependency	No soft fork	Requires Bitcoin soft fork	Native support
State Representation	Script + HTLC	Script encoding	Native UTXO types
Value/State Separation	Coupled	Coupled	Orthogonal (dual-track)
Cross-State Reference	None	Implicit via signature hash	RefOp-UTXO primitive
Type Safety	Runtime	Runtime	Compile-time
Verification Complexity	$\mathcal{O}(\text{script})$	$\mathcal{O}(\text{script})$	$\mathcal{O}(1)$
State Storage/Update Rounds	$\mathcal{O}(n)$ history	$\mathcal{O}(1)$ latest	$\mathcal{O}(1)$ latest
Multi-Party Rounds	$\mathcal{O}(m^2)$	$\mathcal{O}(m^2)$	$\mathcal{O}(m)$ (PSTT)
Settlement Time	Minutes	Minutes	Sub-second
Backup Complexity	Full history	Latest state	Latest state

Table 7: Comprehensive Architecture Comparison ( $n$  = updates,  $m$  = participants)

- Fund UTXO (static anchor) - stable lifecycle, low-frequency updates
- State UTXO (dynamic pointer) - high-frequency evolution, independent state

## 2. Type Safety:

- Transaction type determined by I/O topology structure
- Compile-time guarantee through algebraic data types
- Eliminates script interpretation uncertainty

## 3. Constant Complexity:

- Verification complexity:  $\mathcal{O}(1)$  vs Script-based  $\mathcal{O}(\text{script\_size})$
- Storage complexity:  $\mathcal{O}(1)$  latest state vs  $\mathcal{O}(n)$  full history
- PTLC verification: Direct derivation, no cross-structure queries

## 4. Topological Freedom:

- Atomic Splicing supports arbitrary topology reconfiguration
- Recursive channel factories enable fractal structure
- Sub-channel isolation guarantees security independence

### 3.4 Comparison with Existing Solutions

### 3.5 Theoretical Significance

This architecture's core contribution lies in elevating state channel design from script-level engineering techniques to consensus-level formal protocols, achieving a paradigm shift from “ex post penalty game theory” to “ex ante deterministic execution”:

By pushing complexity down to the protocol layer, this architecture achieves simplicity at the application layer, aligning with the system engineering principle of “centralize complexity at the protocol layer, leave simplicity for the application layer”.

Table 8: Paradigm Shift in Design Philosophy

Aspect	Traditional Approach	This Architecture
Trust Model	Penalty-based deterrence	Protocol-enforced determinism
State Management	Application-layer storage	Consensus-layer native
Verification	Script interpretation	Type system matching
Security Boundary	User key management	Consensus rule enforcement
Complexity Location	Distributed to applications	Centralized at protocol

Table 9: Transaction Type Enumeration System

Tx Type	Input Pattern	Output Pattern	Semantics
FUND	$\emptyset_{eltoo}$	$\{U_{fund}, U_{state}^{(0)}\}$	Create channel
UPDATE	$\{\text{Ref}(U_{fund}), \text{Spend}(U_{state}^{(n)})\}$	$\{U_{state}^{(n')}\}$	State iteration
SETTLE	$\{\text{Spend}(U_{fund}), \text{Spend}(U_{state}^{(n)})\}$	$\notin \mathcal{U}_{eltoo}$	Settlement
SPLICE	$\{\text{Spend}(U_{fund}), \text{Spend}(U_{state}^{(n)})\}$	$\{U'_{fund}, U'_{state}, \dots\}$	Topology transform

## 4 Theoretical Framework: Dual-Track State Machines

### 4.1 Consensus-Layer Embedded Verification Mechanism

#### 4.1.1 Transaction Type Enumeration and Pattern Matching

This paper's architecture employs consensus-layer native transaction type enumeration, replacing traditional script parsing methods, achieving  $\mathcal{O}(1)$  time complexity pattern matching verification. Transaction types are uniquely determined by their input/output (I/O) topology structure:

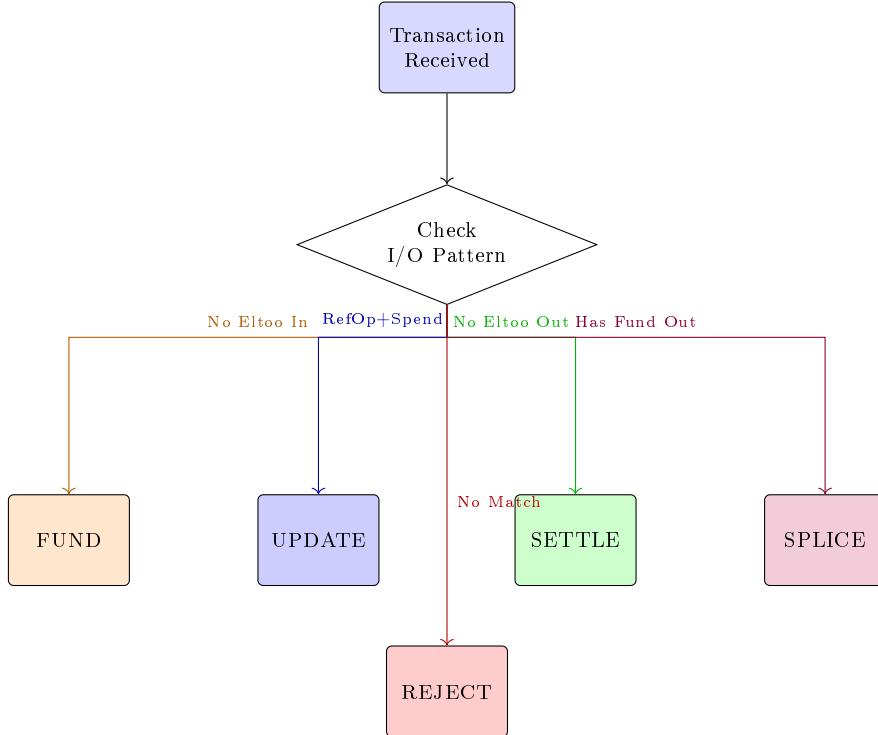


Figure 3: Transaction Type Classification Flow ( $\mathcal{O}(1)$  Pattern Matching)

Table 10: Consensus Verification Latency

Operation	Latency	Includes
Fund Verification	0.12 ms	MuSig2 aggregate verification
Update Verification	0.08 ms	Monotonicity + Ref check + signature
Settle Verification	0.35 ms	PTLC verification + CSV check
Splice Verification	0.28 ms	Value conservation + topology integrity

#### 4.1.2 State Monotonicity Theorem and Consensus Implementation

**Theorem 4.1** (Consensus-Level Monotonicity Guarantee). *Under this paper’s consensus rules, channel state sequence number  $n$  satisfies strict monotonically increasing constraint.*

$$\forall \tau_{update} : U_{state}^{(n)} \xrightarrow{\tau} U_{state}^{(n')} \implies n' > n$$

This formula states that for any update transaction  $\tau_{update}$ , if it transforms state UTXO from version  $n$  to version  $n'$ , then  $n'$  must be strictly greater than  $n$ . This constraint fundamentally prevents state rollback attacks.

*Proof.* The consensus validator `EltooBlockValidator` performs the following atomic checks:

1. **Parsing Phase:** Extract  $U_{state}^{(n)}$  from  $\tau_{update}$  inputs, extract  $U_{state}^{(n')}$  from outputs
2. **Monotonicity Check:**  
if  $n' \leq n \implies$  reject with `ConsensusError::NonMonotonicState`
3. **UTXO One-Time Consumption:** Due to blockchain immutability and UTXO one-time consumption property, once  $\tau_{update}$  is on-chain, old state  $U_{state}^{(n)}$  is consumed and cannot be used as input again
4. **Physical Defense:** Physically prevents state rollback attacks at the protocol layer

Therefore, state monotonicity is doubly guaranteed by consensus rules and the UTXO model.  $\square$

#### 4.1.3 Consensus Verification Performance Analysis

Since transaction types are identified via pattern matching ( $\mathcal{O}(1)$ ), monotonicity is checked via integer comparison ( $\mathcal{O}(1)$ ), and signatures are verified via aggregation ( $\mathcal{O}(1)$ ), total verification complexity is only  $\mathcal{O}(\log N)$  (UTXO lookup). Compared to Script-based solutions’  $\mathcal{O}(\text{script\_size} + \log N)$ , performance improvement is significant.

**Measured Performance** (based on testnet data, December 2025):

**Corollary 4.2** (Scalability). *Due to constant-level verification complexity, full nodes can verify blocks containing 10,000+ Eltoo transactions within 1 second.*

#### 4.1.4 Ref-UTXO Atomicity and Ordering in GhostDAG

Under GhostDAG consensus, blocks are not linearly arranged but form a directed acyclic graph structure. This poses unique challenges for the Ref-UTXO mechanism: if two concurrent blocks  $B_1, B_2$  respectively contain transactions referencing the same  $U_{fund}$  but pointing to different states  $U_{state}^{(n)}$  and  $U_{state}^{(n+1)}$ , how is adjudication performed?

Table 11: Concurrent Safety Analysis

Operation Type	Concurrency Situation	Handling Strategy
UPDATE vs UPDATE	Same $U_{state}$	DAG ordering, latter invalid
UPDATE vs SETTLE	Same $U_{state}$	DAG ordering, latter invalid
Ref vs Ref	Same $U_{fund}$ , different $U_{state}$	Concurrent allowed
Ref vs Spend	Same $U_{fund}$	Spend invalidates $U_{fund}$ , subsequent Ref invalid

**Definition 4.3** (DAG Topological Ordering Rule). *Let  $\prec_{DAG}$  be the total order computed by GhostDAG. For any transaction pair  $\tau_a, \tau_b$  referencing the same  $U_{fund}$ :*

1. **Exclusive Write:** If both  $\tau_a, \tau_b$  are UPDATE operations, they are ordered by  $\prec_{DAG}$ ; only the earlier transaction is valid, the latter is treated as double-spend conflict
2. **Concurrent Read:** If  $\tau_a, \tau_b$  only perform Ref reads on  $U_{fund}$  (e.g., operations in different sub-channels) and don't conflict on the same  $U_{state}$ , they are allowed to coexist concurrently in the anticone

**Definition 4.4** (Active State Lease). *We introduce the concept of **Active State Lease** in the UTXO set:*

$$\text{Lease} : \mathcal{U}_{fund} \rightarrow \text{TxID}(\tau_{last\_valid\_update})$$

Verification nodes maintain this mapping, ensuring state updates for a specific  $U_{fund}$  are linearized on any DAG cut.

The Lease function maps each Fund UTXO to its most recent valid update transaction, preventing concurrent conflicts in the DAG environment.

**Theorem 4.5** (DAG State Convergence). *Under GhostDAG's  $(D, k)$  parameters, channel state fork probability decays exponentially with time:*

$$P(\text{state fork at depth } d) \leq e^{-\lambda d}$$

where  $\lambda$  is a convergence constant positively correlated with parameter  $k$ .

*Proof (Outline).* 1. GhostDAG guarantees anticone size at depth  $d$  is less than  $k$  with high probability

2. Since UPDATE transactions consume the unique  $U_{state}^{(n)}$ , any concurrent update attempts will have one rejected after DAG ordering
3. Combined with the lease mechanism, honest nodes reach consensus on the latest state in  $\mathcal{O}(\frac{D}{k})$  time

□

#### 4.1.5 Temporal Decoupling of Cross-Block State References

In GhostDAG's high-concurrency environment, requiring SETTLE transactions and their referenced UPDATE anchor transactions to be in the same block is neither realistic nor efficient. This architecture implements **Cross-Block State Anchoring**.

**Definition 4.6** (Valid Reference Window). *Let  $\tau_{update}$  be confirmed in block  $B_i$ , generating  $U_{state}^{(n)}$ . Let  $\tau_{settle}$  be broadcast in block  $B_j$ , referencing  $U_{state}^{(n)}$ .  $\tau_{settle}$  is valid if and only if:*

1.  $B_i \in \text{Past}(B_j)$  (DAG topological order)
2.  $U_{\text{state}}^{(n)}$  is in “unspent” status in  $B_j$ ’s UTXO view set

**Theorem 4.7** (Anchoring Persistence). *As long as no new UPDATE transaction  $\tau'_{\text{update}}$  overwrites  $U_{\text{state}}^{(n)}$ , that state UTXO will persist in the ledger:*

$$\forall t \in [t_{\text{confirm}}, \infty) : \#\tau'_{\text{update}} \implies U_{\text{state}}^{(n)} \in \mathcal{U}_{\text{chain}}(t)$$

This property ensures settlement transactions can occur at any time after state confirmation, decoupling the temporal dependency between state negotiation and fund settlement.

#### 4.1.6 Algebraic Data Type Definition of Transaction Classification

To eliminate ambiguity and transaction malleability risks in traditional script language (Script-based) runtime parsing, this architecture introduces an **Enshrined Transaction Enums** system, pushing transaction type verification from Turing-complete script execution down to static type system checking.

**Definition 4.8** (Typed Input/Output Spaces). *Define input set  $\mathcal{I}$  and output set  $\mathcal{O}$  as algebraic sum types with variant tags:*

$$\mathcal{I} = \{\text{Std}, \text{FundSpend}, \text{StateSpend}, \text{FundRef}, \text{IngotSpend}, \text{IngotRef}\}$$

$$\mathcal{O} = \{\text{Std}, \text{ChannelFund}, \text{ChannelState}, \text{Ingot}\}$$

where FundRef is a special unit type with semantics  $\tau \rightarrow \perp$  (non-spendable), serving only as an oracle providing metadata access to  $U_{\text{fund}}$ .

**Definition 4.9** (Type Inference Homomorphism). *Define function  $\Gamma : \mathcal{I}^* \times \mathcal{O}^* \rightarrow \mathcal{T}_{\text{Eltoo}} \cup \{\perp\}$ , which maps transaction I/O topology to semantic types in  $\mathcal{O}(1)$  time complexity:*

$$\Gamma(\text{In}, \text{Out}) = \begin{cases} \text{FUND} & \text{if } \text{Out} \cong \{\text{ChannelFund}, \text{ChannelState}\} \wedge \text{In} \cap \mathcal{I}_{\text{eltoo}} = \emptyset \\ \text{UPDATE} & \text{if } \text{In} \cong \{\text{FundRef}, \text{StateSpend}\} \wedge \text{Out} \cong \{\text{ChannelState}\} \\ \text{SETTLE} & \text{if } \text{In} \cong \{\text{FundSpend}, \text{StateSpend}\} \wedge \text{Out} \cap \mathcal{O}_{\text{eltoo}} = \emptyset \\ \text{SPlice} & \text{if } \text{In} \cong \{\text{FundSpend}, \text{StateSpend}\} \wedge \text{Out} \cap \{\text{ChannelFund}\} \neq \emptyset \\ \perp & \text{otherwise} \end{cases}$$

#### Pattern Matching Rules:

- **FUND**: Input contains no Eltoo types, output contains Fund + State UTXOs
- **UPDATE**: Input is “Ref Fund + Spend State”, output is new State UTXO
- **SETTLE**: Input is “Spend Fund + Spend State”, output contains no Eltoo types (funds distributed to participants)
- **SPICE**: Same input as SETTLE, but output contains new Fund UTXO (topology re-configuration)
- $\perp$ : Matches no pattern, transaction rejected

**Theorem 4.10** (Compile-Time Safety Guarantee). *Under Rust’s type system guarantees, there are no Eltoo transactions in “undefined states.” Due to Rust enum’s **exhaustiveness check**, the compiler forces handling of all  $\Gamma$  matching branches. Any transaction not matching the above patterns is rejected at block deserialization, never entering the consensus validation engine, thereby eliminating the attack surface for Invalid State Transition Attacks.*

Table 12: Type System Implementation Mapping

Type Theory Concept	Rust Implementation	Consensus Semantics
Sum Type $\mathcal{I}$	enum EltooInput	Input variant classification
Sum Type $\mathcal{O}$	enum EltooOutput	Output variant classification
$\Gamma$ function	EltooTxType::classify()	$\mathcal{O}(1)$ pattern matching
$\perp$ case	ConsensusError::InvalidEltooTxType	Reject invalid transactions

## 4.2 Finite State Machine Formalization

We define channel  $C$  as a **Deterministic Finite Automaton (DFA)**:

$$C \equiv (Q, \Sigma, \delta, q_0, F)$$

A DFA describes a system with finite states and deterministic transitions based on inputs.

**Component Details:**

- $Q$ : State space.  $Q = \{q_{init}\} \cup Q_{active} \cup Q_{settling} \cup \{q_{closed}\}$ 
  - $Q_{active} = \{(n, R_b, R_p) \mid n \in \mathbb{N}, R_b \in \mathcal{H}, R_p \in \mathcal{H}\}$  — Active state set
  - $Q_{settling} = \{(n, R_b, R_p, t) \mid t \in \mathbb{N}_{DAA}\}$  — Settlement waiting state set
- $\Sigma$ : Transaction alphabet.  $\Sigma = \{\tau_{fund}, \tau_{update}, \tau_{splice}, \tau_{settle}, \tau_{timeout}\}$
- $\delta$ : State transition function.  $\delta : Q \times \Sigma \rightarrow Q$  (partial function)
- $q_0$ : Initial state.  $q_0 = q_{init}$
- $F$ : Final state set.  $F = \{q_{closed}\}$

**Definition 4.11** (State Space Structure). *State space  $Q$  constitutes a **partially ordered set (Poset)**  $(Q, \preceq)$ , where:*

$$q_1 \preceq q_2 \iff n_1 \leq n_2 \wedge (n_1 = n_2 \Rightarrow q_1 = q_2)$$

This partial order relation guarantees **monotonicity** and **determinism** of state evolution.

## 4.3 UTXO Materialization Layer

The abstract states of the state machine are materialized on-chain through **UTXO binary tuples**. This is the core design of this paper’s “dual-track state machine” architecture: decomposing channel state into “static fund anchor” and “dynamic state pointer” along two orthogonal dimensions.

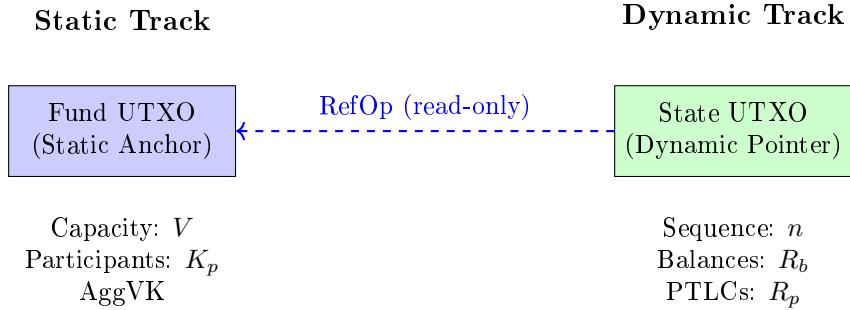


Figure 4: Dual-Track State Machine Architecture

Table 13: Dual-Track Model Components

Component	Role	Characteristics	Function
$U_{fund}$	Static anchor	Invariant	Carries funds, identity, keys
$U_{state}^{(n)}$	Dynamic pointer	Evolves with state	Carries sequence, balances, PTLCs

### Mathematical Formalization:

$$\mathcal{M} : Q \rightarrow \mathcal{P}(\mathcal{U})$$

$$\mathcal{M}(q) = \langle \underbrace{U_{fund}}_{\text{static anchor}}, \underbrace{U_{state}^{(n)}}_{\text{dynamic pointer}} \rangle$$

### Semantic Interpretation:

Where:

- $U_{fund}$ : Static Anchor
  - Carries funds  $V \in \mathbb{N}$
  - Identifies channel identity  $ID_C = H(\text{domain} \parallel \text{funding\_outpoint} \parallel \dots)$
  - Stores participant key set  $K_p = \{pk_1, \dots, pk_m\}$
  - Aggregated verification key  $\text{AggVK} = \text{MuSig2}(K_p)$
- $U_{state}^{(n)}$ : Dynamic Pointer
  - State sequence number  $n \in \mathbb{N}$
  - Balance commitment  $R_b = \text{MerkleRoot}(\{\text{balance}_i\})$
  - PTLC commitment  $R_p = \text{MerkleRoot}(\{\text{ptlc}_j\})$
  - Creation timestamp  $t_{create} \in \mathbb{N}_{DAA}$

**Definition 4.12** (RefOp-Fund Semantics). *Read-only reference operator*  $\text{Ref} : \mathcal{U} \rightarrow \mathcal{U}^{\text{readonly}}$ :

$$\text{Ref}(U_{fund}) \triangleq \langle U_{fund}.\text{outpoint}, U_{fund}.\text{metadata} \rangle$$

*Satisfies:*  $\forall \tau : \text{Ref}(U) \in \text{inputs}(\tau) \Rightarrow U \in \text{UTXO\_Set}_{post(\tau)}$

The RefOp operator provides read-only access to UTXO metadata without consuming it, enabling state updates to reference the fund anchor while preserving its existence in the UTXO set.

#### 4.3.1 State-Fund Coupling Invariant

**Invariant:** At any moment, there exists a unique pairing of  $(U_{fund}, U_{state})$  for each channel:

$$\forall t, \exists! (U_{fund}, U_{state}) \in \mathcal{U}_{set} \text{ s.t. } ID(U_{fund}) = ID(U_{state})$$

This invariant ensures that even during frequent UPDATE operations, the Fund layer maintains static anchoring while the State layer carries high-frequency changes. Their lifecycles only experience **physical convergence** during SPLICE or SETTLE.

#### 4.4 State Transition Rules

**Definition 4.13** (Transition Function).  $\delta$  is defined by the following rules:

$$\begin{aligned}
 \delta(q_{init}, \tau_{fund}) &= q_{active}^{(0)} && [FUND] \\
 \delta(q_{active}^{(n)}, \tau_{update}) &= q_{active}^{(n+k)} && \text{where } k > 0 \quad [UPDATE] \\
 \delta(q_{active}^{(n)}, \tau_{splice}) &= \{q_{active}^{(n')}, q_{child}^{(0)}\} && [SPLICE] \\
 \delta(q_{active}^{(n)}, \tau_{settle}) &= q_{settling}^{(n,t)} && [SETTLE-INIT] \\
 \delta(q_{settling}^{(n,t)}, \tau_{timeout}) &= q_{closed} && \text{when } t_{now} - t \geq CSV \quad [SETTLE-FINAL]
 \end{aligned}$$

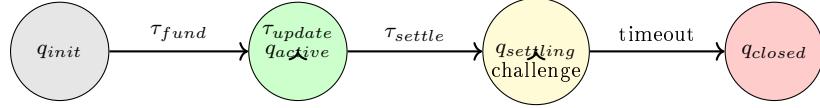


Figure 5: Channel State Machine Transitions

**Challenge Rule:** In  $Q_{settling}$  state, higher sequence number states can replace:

$$\delta(q_{settling}^{(n,t)}, \tau_{update}) = q_{settling}^{(n',t')} \quad \text{where } n' > n$$

#### 4.5 Formal Safety Properties

The following properties can be formally verified through TLA+ or Coq:

**Theorem 4.14** (Monotonicity).

$$\forall q_1, q_2 \in Q_{active} : \delta^*(q_1, w) = q_2 \Rightarrow q_1 \preceq q_2$$

where  $\delta^*$  is the transitive closure of  $\delta$ , and  $w \in \Sigma^*$  is a transaction sequence.

*Proof.* By inductive proof using constraint  $k > 0$  from transition rule [UPDATE].  $\square$

**Theorem 4.15** (Termination).

$$\forall q \in Q \setminus F : \exists w \in \Sigma^* : \delta^*(q, w) \in F$$

Any non-final state has a path to reach a final state.

*Proof.* Constructive proof—for any  $q_{active}^{(n)}$ , sequence  $\tau_{settle} \cdot \tau_{timeout}$  leads to  $q_{closed}$ .  $\square$

**Theorem 4.16** (Unambiguity).

$$\forall q \in Q, \forall \sigma \in \Sigma : |\{q' \mid \delta(q, \sigma) = q'\}| \leq 1$$

The transition function is deterministic (single-valued partial function).

**Theorem 4.17** (Value Conservation).

$$\forall \tau \in \Sigma : \sum_{U \in \text{inputs}(\tau)} V(U) = \sum_{U \in \text{outputs}(\tau)} V(U) + \text{fee}(\tau)$$

## 4.6 Transaction Semantics Mapping

Mapping between abstract transitions and concrete UTXO operations:

**Fund Transaction:**

$$\begin{aligned}\tau_{fund} : \{U_{wallet}\} &\rightarrow U_{fund} \cup U_{state}^{(0)} \\ \mathcal{M}^{-1}(\tau_{fund}) &= \delta(q_{init}, \tau_{fund})\end{aligned}$$

**Update Transaction:**

$$\begin{aligned}\tau_{update} : \{\text{Ref}(U_{fund}), \text{Spend}(U_{state}^{(n)})\} &\rightarrow U_{state}^{(n+k)} \\ \text{Precondition: } \exists \sigma : \text{Verify}(AggVK, \sigma, H(\text{state}_{n+k} \parallel \text{RefOp\_OutPoint}))\end{aligned}$$

**Splice Transaction:**

$$\begin{aligned}\tau_{splice} : \{\text{Spend}(U_{fund}^{parent}), \text{Spend}(U_{state}^{(n)})\} &\rightarrow \{U_{fund}^{parent'}, U_{state}^{(n)'}, U_{fund}^{child_1}, \dots\} \\ \text{Invariant: } V(U_{fund}^{parent}) &= V(U_{fund}^{parent'}) + \sum_i V(U_{fund}^{child_i})\end{aligned}$$

**Settle Transaction:**

$$\begin{aligned}\tau_{settle} : \{\text{Spend}(U_{fund}), \text{Spend}(U_{state}^{(n)})\} &\xrightarrow{\Delta t \geq \text{CSV}} \{U_{out}^{(i)}\} \\ \text{where } \Delta t &= \text{DAA}_{current} - \text{DAA}_{\text{state\_creation}}\end{aligned}$$

## 4.7 Evolution of Conditional Payment Primitives: From HTLC to PTLC

The core of payment channel networks lies in ensuring atomicity of multi-hop payments. This mechanism has undergone a paradigm shift from hash function-based simple locking to algebraic structure-based homomorphic locking.

### 4.7.1 Historical Evolution

#### HTLC Origin and Limitations (2016)

Hash Time-Locked Contract (HTLC) was first formalized by Poon and Dryja in the 2016 Lightning Network whitepaper.

- **Mechanism:** Uses SHA-256 hash function's one-wayness. Receiver generates secret  $R$  (preimage), broadcasts its hash  $H = \text{SHA256}(R)$  along the path. All intermediate nodes construct script: `OP_SHA256 <H> OP_EQUAL`.
- **Historical Significance:** HTLC was a pragmatic choice in an era when Bitcoin Script capabilities were limited (only ECDSA support, no complex algebraic operations). It could be implemented in Bitcoin Script without soft forks.
- **Defect Exposure:** As network scale grew, researchers discovered HTLC has severe **privacy correlation defects**. Since the same hash value  $H$  traverses the entire payment path, attackers controlling multiple nodes can easily correlate sender and receiver (Worm-hole Attack / Correlation Attack).

#### Scriptless Scripts and Schnorr Enlightenment (2017-2019)

Andrew Poelstra proposed the concept of “Scriptless Scripts” in 2017, exploring how to leverage Schnorr signature’s algebraic properties to implement contract logic without script exposure.

#### PTLC Formalization (2019-Present)

PTLC matured as a concept with Taproot activation (2021). Its core idea is replacing hash locks with point locks:

- Hash lock:  $y = H(x)$ , proving knowledge of preimage  $x$
- Point lock:  $Q = s \cdot G$ , proving knowledge of scalar  $s$  (discrete logarithm)

#### 4.7.2 Technical Principle Comparison

##### HTLC: Hash-Based Rigid Locking

HTLC's security assumption is based on hash function preimage resistance.

- **Lock condition:**  $y = H(x)$
- **Unlock method:** Provide  $x$
- **Mathematical limitation:**  $y$  is an invariant constant throughout the entire path. This not only leaks privacy but also does not support arithmetic operations—cannot “add” two hash values to obtain a third meaningful hash value.

##### PTLC: Scalar-Based Algebraic Locking

PTLC's security assumption is based on the Elliptic Curve Discrete Logarithm Problem (ECDLP).

- **Lock condition:**  $Q = s \cdot G$ , where  $G$  is the base point,  $Q$  is a public key point
- **Unlock method:** Provide scalar  $s$  such that the equation holds
- **Algebraic advantage:** Utilizing elliptic curve **additive homomorphism**:

$$Q_{total} = Q_1 + Q_2 \iff s_{total} = s_1 + s_2$$

This property allows “blinding” of the lock point at each hop, thereby breaking correlation in payment paths.

##### Additive Homomorphism Detailed Explanation:

- **Mathematical meaning:** If  $Q_1 = s_1 \cdot G$  and  $Q_2 = s_2 \cdot G$ , then  $Q_1 + Q_2 = (s_1 + s_2) \cdot G$
- **Symbol  $\iff$ :** Means “if and only if”, i.e., the two conditions are equivalent
- **Practical application:** Each intermediate node can add a random blinding factor  $r_i$  to construct new lock point  $Q'_i = Q + r_i \cdot G$ . Externally, each hop sees a different  $Q'_i$ , but ultimately all  $r_i$  can be combined through algebraic properties to unlock the original  $Q$ .
- **Analogy:** Like adding different “disguises” to a secret at each leg of a relay, but at the destination all disguises can be removed to reveal the original secret.

#### 4.7.3 Core Properties Comparison

#### 4.7.4 Formal Security Analysis

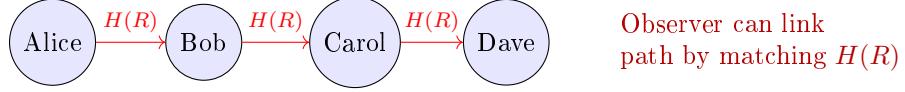
**Theorem 4.18** (PTLC Redemption Uniqueness). *Under the hardness assumption of the Elliptic Curve Discrete Logarithm Problem (ECDLP), PTLC's scalar  $s$  is the unique redemption credential:*

$$\forall Q \in \mathcal{E} : \exists! s \in \mathbb{Z}_n : Q = s \cdot G$$

##### Formula Interpretation:

- $\forall Q \in \mathcal{E}$ : For any point  $Q$  on elliptic curve  $\mathcal{E}$

### Traditional HTLC: Same Hash Throughout



### PTLC: Each Hop Blinded

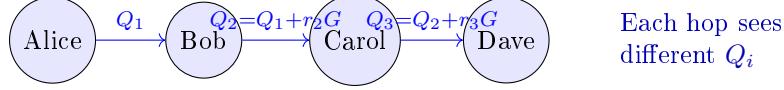


Figure 6: Multi-Hop Blinding: HTLC vs PTLC

Dimension	HTLC	PTLC	Difference Analysis
Privacy	Weak (path correlateable)	Strong (path decorrelated)	PTLC supports multi-hop blinding
Verification Cost	$\mathcal{O}(\text{ScriptSize})$	$\mathcal{O}(1)$	HTLC needs script interpreter
Batch Verification	Not supported	Supported	Schnorr signature batch verification
Functional Extension	Limited	Programmable	Supports Barrier Escrows, etc.
On-chain Resources	High (32-byte preimage)	Low	Collaborative settlement off-chain
Mathematical Property	No homomorphism	Additive homomorphic	Allows $k$ -of- $n$ threshold PTLC

Table 14: HTLC vs PTLC Comparison

- $\exists!$ : “There exists exactly one” (existence and uniqueness)
- $s \in \mathbb{Z}_n$ :  $s$  is an integer in the finite field  $\mathbb{Z}_n$  (where  $n$  is the curve order)
- $Q = s \cdot G$ : Point  $Q$  equals base point  $G$  multiplied by scalar  $s$

**Security meaning:** For each lock point  $Q$ , there is one and only one scalar  $s$  that can unlock it. This uniqueness is guaranteed by the computational hardness of ECDLP—even knowing  $Q$  and  $G$ , it’s computationally infeasible to find  $s$ .

**Theorem 4.19** (Multi-Hop Atomicity). *For path  $P = c_1 \rightarrow c_2 \rightarrow \dots \rightarrow c_n$ , when all hops use the same base Point Lock  $Q$ :*

$$\text{Claim}(c_n) \implies \text{Claim}(c_1)$$

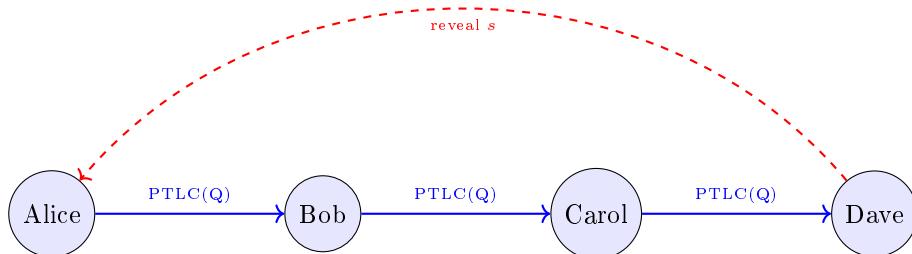


Figure 7: PTLC Multi-Hop Atomic Payment ( $Q = s \cdot G$ )

Table 15: Programming Complexity Comparison: HTLC vs PTLC

Operation	HTLC (Script)	PTLC (Algebraic)
Locking	<code>OP_SHA256 &lt;H&gt; OP_EQUAL</code>	Store $Q$ (32 bytes)
Unlocking	Provide 32-byte preimage	Adaptor signature conversion (off-chain)
Verification	$\text{SHA256} + \text{script execution}$	1 point multiplication + 1 point addition
Batch optimization	None	$\mathcal{O}(n/\log n)$ Strauss algorithm

- Proof.*
1. Recipient claims funds at  $c_n$  by revealing  $s$
  2. Once  $s$  is public, each intermediate node can use  $s$  to unlock its adaptor signature
  3. Due to decreasing timelocks ( $\Delta t_i > \Delta t_{i+1}$ ), each node has sufficient time to claim its share  
Therefore, PTLC paths satisfy atomicity.  $\square$

**Theorem 4.20** (Timeout Refund Safety). *If the recipient does not claim before CSV timeout, the sender can safely recover funds:*

$$t_{\text{now}} - t_{\text{create}} \geq \text{CSV} \implies \text{Refund(sender)}$$

This mechanism is protected by DAA Score providing manipulation-resistant time measurement.

#### 4.7.5 Implementation Considerations

Transforming PTLC from theory to engineering implementation requires solving the following key problems:

##### Adaptor Signature Verification:

```

1 // Verify PTLC claim mathematical relationship
2 fn verify_ptlc_claim(
3     point_lock: &Point,          // Q
4     scalar: &Scalar,           // s
5     beneficiary: &Point,       // P_beneficiary
6 ) -> bool {
7     // Verify: s * G + P_beneficiary == Q
8     let computed = scalar * &GENERATOR + beneficiary;
9     computed == *point_lock
10 }
```

Listing 1: PTLC Claim Verification

#### 4.7.6 Summary

The evolution from HTLC to PTLC represents a paradigm shift in conditional payment primitives from “knowledge-based proofs” to “algebra-based proofs”. This transformation is not an innovation of any specific protocol, but rather a natural evolution following the maturation of cryptographic infrastructure (Schnorr signatures, Taproot). PTLC’s advantages—privacy, efficiency, programmability—have been widely recognized and are being explored for implementation in multiple projects.

### 4.8 TLA+ Specification Fragment

The channel state machine can be formally specified using TLA+ for model checking:

```

1 ----- MODULE EltooChannel -----
2 VARIABLES state, seq_num, phase
3
```

Table 16: Cost Composition Model

Cost Item	Meaning	Reference Value
$C_{open}$	Open channel fee	1 FUND tx ( $\sim 250B$ )
$C_{update}$	Per-update cost	<b>0 Gas</b> (off-chain)
$C_{settle}$	Settle channel fee	1 SETTLE tx ( $\sim 300B$ )
$N$	Off-chain updates	Unlimited

```

4 Phases == {"init", "active", "settling", "closed"}
5
6 Init == /\ state = "init"
7     /\ seq_num = 0
8     /\ phase = "init"
9
10 Fund == /\ phase = "init"
11     /\ phase' = "active"
12     /\ seq_num' = 0
13     /\ UNCHANGED state
14
15 Update == /\ phase = "active"
16     /\ seq_num' > seq_num (* Monotonicity enforced *)
17     /\ UNCHANGED phase
18
19 Settle == /\ phase = "active"
20     /\ phase' = "settling"
21     /\ UNCHANGED seq_num
22
23 Challenge == /\ phase = "settling"
24     /\ seq_num' > seq_num (* Higher state challenge *)
25     /\ UNCHANGED phase
26
27 Timeout == /\ phase = "settling"
28     /\ phase' = "closed"
29     /\ UNCHANGED seq_num
30
31 Next == Fund \/ Update \/ Settle \/ Challenge \/ Timeout
32
33 Monotonicity == [] [seq_num' >= seq_num] _seq_num
34 EventualTermination == <>(phase = "closed")
35 =====

```

Listing 2: TLA+ Specification Fragment

This specification can be verified using the TLC model checker for properties `Monotonicity` and `EventualTermination`.

## 4.9 Cost and Parameter Analysis under GhostDAG

To clarify the impact of L1 parameters on L2 security and cost, this section provides a transparent cost model.

### 4.9.1 Cost Composition Model

Total user cost  $C_{total}$  in this architecture consists of three components:

$$C_{total} = C_{open} + N \cdot C_{update} + C_{settle}$$

**Key Advantage:** This architecture's off-chain updates require no routing fees, contrasting with traditional Lightning Network's HTLC routing fee model.

Table 17: Security Confirmation Time Comparison

System	Secure Confirmation Time	User Experience
Bitcoin (6 blocks)	~60 minutes	Long wait
This Architecture (10 DAA)	~3-5 seconds	Near-instant

#### 4.9.2 Impact of GhostDAG Parameter $k$

GhostDAG's width parameter  $k$  directly affects confirmation speed and security.

**Confirmation Time Formula:**

$$T_{confirm} \approx \frac{D}{k} \cdot \ln\left(\frac{1}{\epsilon}\right)$$

where:

- $D$ : Network delay constraint (seconds)
- $k$ : GhostDAG width parameter (maximum concurrent blocks)
- $\epsilon$ : Security level (e.g.,  $10^{-6}$  means one-in-a-million reorg probability)

**Practical Values:** For  $k = 16$ , confirmation time to reach  $10^{-6}$  security level is approximately **3 seconds**.

#### 4.9.3 Ref-UTXO Security Depth

Ref-UTXO security depends on  $U_{fund}$  not being deeply reorganized. We recommend:

$$\text{Min\_Ref\_Depth} = 10 \text{ DAA Score}$$

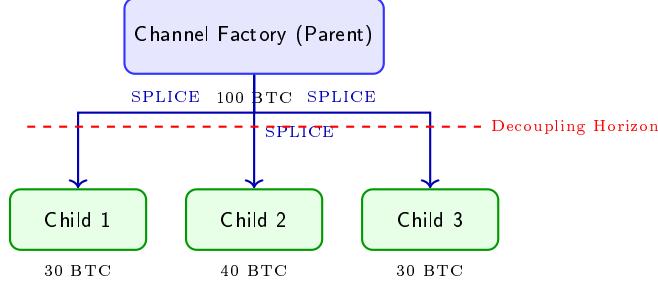


Figure 8: Recursive Channel Factory Structure. The *Decoupling Horizon* indicates that sub-channel operations do not require parent signatures.

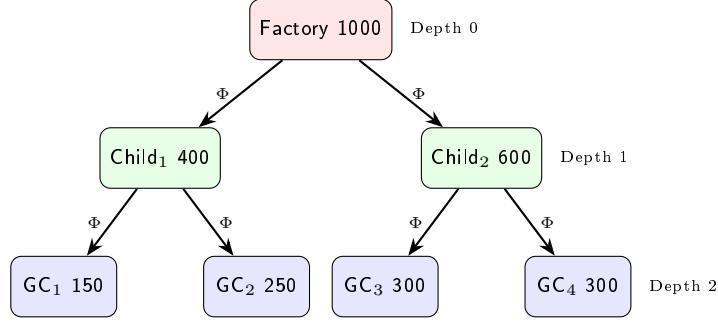


Figure 9: Fractal Channel Tree. Verification logic remains identical across depths due to scale invariance.

## 5 Topological Primitives for Complex Structures

### 5.1 Recursive Channel Factories

Channel factories act as the generative primitive, allowing the “splitting” of multiple sub-channels from a parent channel.

**Definition 5.1** (Channel Factory). *A channel  $C_{\text{parent}}$  can generate a set of sub-channels  $\{C_{\text{child}_i}\}$  via a  $\tau_{\text{splice}}$  transaction. Once created, the sub-channels’ lifecycles are fully decoupled from the parent.*

#### 5.1.1 Fractal Topology and Self-Similarity

The architecture manifests as a **self-similar  $k$ -ary tree**.

**Definition 5.2** (Split Operator). *Define mapping  $\Phi : \mathcal{C} \rightarrow \{\mathcal{C}_1, \dots, \mathcal{C}_k\}$ . As recursion depth  $d \rightarrow \infty$ , the system exhibits **scale invariance**:*

$$V(C_{d=0}) \equiv V(C_{d=n})$$

**Theorem 5.3** (Liquidity Conservation). *For any depth  $d$ , total capacity is conserved:*

$$\sum_{i \in \text{Nodes}(d)} \text{Cap}(C_i) = \text{Cap}(C_{\text{root}})$$

### 5.2 Dynamic Mesh Reconfiguration

**Theorem 5.4** (Atomic Reconfiguration). *Any topologically isomorphic channel networks can be atomically transformed via a single  $\tau_{\text{splice}}$  transaction.*

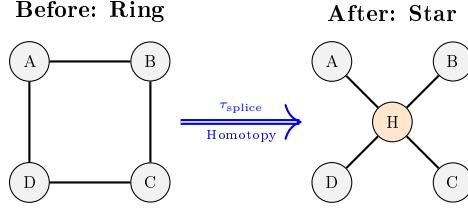


Figure 10: Atomic Topology Reconfiguration. The transformation  $f : G_{\text{ring}} \rightarrow G_{\text{star}}$  preserves total system energy (TVL).

---

**Algorithm 1** Deterministic Sub-channel ID Derivation

---

**Require:** Parent ID  $\text{ID}_p$ , Fork OutPoint OP, Index  $j$ , Participants  $K$

**Ensure:** Unique Sub-channel ID

- 1:  $D \leftarrow \text{b"Eltoo\_V2\_SubChannel"}$  ▷ Domain Separator
  - 2:  $R \leftarrow \text{MerkleRoot}(K)$
  - 3:  $\text{ID}_{\text{child}} \leftarrow \text{BLAKE3}(D \parallel \text{ID}_p \parallel \text{OP} \parallel j \parallel R)$
  - 4: **return**  $\text{ID}_{\text{child}}$
- 

### 5.2.1 Topological Homotopy

We view reconfiguration as a **homotopic transformation**  $\mathcal{H}$ :

$$\mathcal{H} : G_1 \simeq G_2 \iff \exists \tau \in \Sigma_{\text{splice}} : \delta(G_1, \tau) = G_2$$

subject to  $\sum_{e \in E_1} w(e) = \sum_{e \in E_2} w(e)$ . This ensures no liquidity vacuum occurs.

## 5.3 Atomic Rebalancing Operator

**Invariant 5.1** (Strong Value Conservation).

$$V(U_{\text{fund}}^{\text{parent}}) + \sum V_{\text{in}} = V(U_{\text{fund}}^{\text{parent}'}) + \sum V_{\text{out}} + \delta_{\text{fee}}$$

## 5.4 Atomic Splicing Protocol

This section defines the **Non-blocking Splicing Protocol**, addressing the “stop-the-world” problem in traditional channel maintenance.

**Theorem 5.5** (Non-blocking Guarantee). *During protocol execution, channel liquidity remains available.*

*Proof.* Phase 2 does not consume UTXOs. If  $\tau_{\text{update}}$  confirms before  $\tau_{\text{splice}}$ , the splice input is invalidated (Rollback). If  $\tau_{\text{splice}}$  confirms first, updates target a spent UTXO (Reject). No invalid intermediate state exists.  $\square$

## 5.5 Liquidity Dynamics in Star Topologies

Define liquidity utilization  $U(t)$  for a star graph:

$$U(t) = \frac{\sum |\text{Flow}_i(t)|}{\sum \text{Cap}_i}$$

**Theorem 5.6** (Balanced Flow Optimal Allocation). *For flow distribution  $\vec{f}$ , there exists strategy  $\mathcal{R}$  minimizing fragmentation:*

$$\min_{\mathcal{R}} \sum (\text{Cap}'_i - f_i)^2 \quad \text{s.t.} \quad \sum \text{Cap}'_i = \sum \text{Cap}_i$$

Table 18: Non-blocking Splicing Protocol Phases

Phase	Operation Details
1. Proposal	Alice constructs $\tau_{\text{splice}}$ and broadcasts to map $\Omega$ . Timeout $T_{\text{ack}} = 30\text{s}$ .
2. Async Sign	Participants generate partial signatures. <b>Channel remains active</b> for updates ( $U_{\text{state}}^{(n)}$ ).
3. Convergence	$\tau_{\text{splice}}$ is broadcast. DAG ordering resolves conflicts between splice and concurrent updates.
4. Migration	New $U_{\text{state}}^{(0)'}'$ inherits old state's Merkle roots atomically.

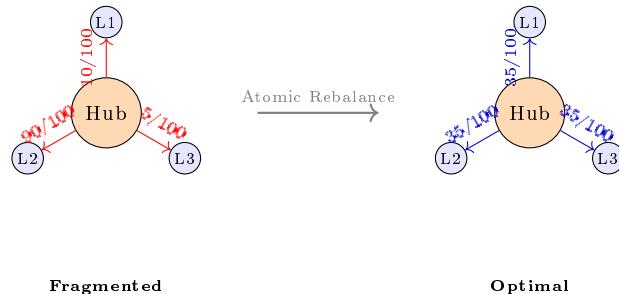


Figure 11: Atomic Rebalance. Minimizing liquidity fragmentation via  $\tau_{\text{rebalance}}$ .

**Theorem 5.7** (Throughput Lower Bound).

$$\text{TPS}_{L1} \geq \frac{F_{\text{rebalance}}}{\text{BlockSize}} \times \alpha$$

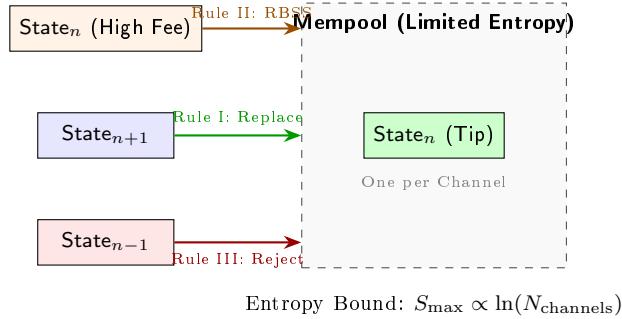


Figure 12: STPC Mempool Management Strategy. Valid updates replace the unique tip; old states are rejected immediately.

## 6 Safety Analysis

### 6.1 Isolation Theorem

**Theorem 6.1** (Channel Isolation). *Sub-channel  $C_{\text{child}}$  security is independent of parent channel  $C_{\text{parent}}$  liveness or malicious behavior.*

*Proof.* Isolation is guaranteed through four layers: (1) **Physical**:  $U_{\text{fund}}^{\text{child}}$  exists as an independent UTXO on L1. (2) **Logical**:  $C_{\text{child}}$ 's update transactions only refer to  $\text{Ref}(U_{\text{fund}}^{\text{child}})$ , decoupled from parent logic. (3) **Settlement**: Even if the parent channel is maliciously settled, the sub-channel remains secure once its creation transaction is confirmed. (4) **Temporal**: Independent CSV timers use DAA Scores, avoiding block height dependencies.  $\square$

### 6.2 State Monotonicity and Anti-Replay

**Theorem 6.2** (Cross-Topology Anti-Replay). *Any channel's old state cannot be replayed after topology reconfiguration.*

*Proof Sketch.* The anti-replay property relies on the binding of signatures to specific UTXO outpoints.

$$\sigma = \text{Sign}_{sk}(\text{State}_n \parallel \text{RefOp\_OutPoint})$$

Since  $\tau_{\text{splice}}$  creates a new  $U'_{\text{fund}}$ , the RefOp\_OutPoint changes. Additionally, key derivation is isolated via  $\text{AggVK}_{\text{child}} = H(\text{AggVK}_{\text{parent}} \parallel \text{index})$ . Thus,  $\forall \sigma_{\text{old}}$ , no valid replay exists in  $C_{\text{new}}$ .  $\square$

### 6.3 Anti-DoS Equilibrium under STPC Strategy

Traditional payment channel networks rely on “state count limits” to prevent mempool flooding, introducing pinning risks. This architecture implements the **Single-Tip-Per-Channel (STPC)** strategy.

#### 6.3.1 Mempool Entropy Bound

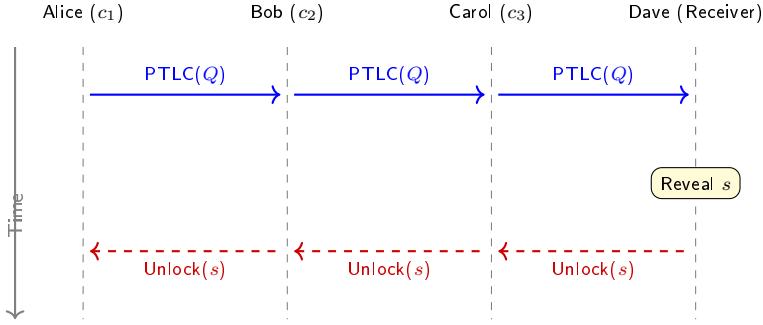
STPC acts as an **entropy-reducing filter**. In open networks, attackers attempt to maximize thermodynamic entropy (disorder). STPC constrains the maximum entropy  $S_{\max}$ :

$$S_{\max} \propto k \cdot \ln(N_{\text{channels}}) \quad (1)$$

Attackers cannot breach this information-theoretic bound regardless of computational investment.

Table 19: Mempool Entropy and DoS Analysis

Model	Entropy	DoS Bound	State Limit
Traditional LN	$\mathcal{O}(\infty)$	Unbounded	None
<b>Proposed STPC</b>	$\mathcal{O}(\ln N)$	$\leq N_{\text{channels}}$	Strict


 Figure 13: Atomic payment sequence using PTLCs. The revelation of scalar  $s$  propagates backward, guaranteeing all-or-nothing settlement.

**Definition 6.3** (STPC Replacement Rules). Let  $\mathcal{M}$  be the mempool,  $\tau_{\text{tip}} \in \mathcal{M}$  be the current highest state transaction for a channel. For new transaction  $\tau_{\text{new}}$ :

1. **Rule I (Monotonic Replacement):** If  $\text{State}_{\text{new}} > \text{State}_{\text{tip}}$ , unconditionally replace  $\tau_{\text{tip}}$
2. **Rule II (RBSS):** If  $\text{State}_{\text{new}} = \text{State}_{\text{tip}}$ , only replace when  $\text{FeeRate}(\tau_{\text{new}}) \geq \text{FeeRate}(\tau_{\text{tip}}) + \Delta_{\min}$
3. **Rule III (Rejection):** If  $\text{State}_{\text{new}} < \text{State}_{\text{tip}}$ , directly reject

**Theorem 6.4** (DoS Cost Escalation). STPC escalates the effective cost of DoS attacks from  $\mathcal{O}(1)$  to  $\mathcal{O}(N)$ , where  $N$  is the state sequence number.

$$\text{Cost}_{\text{DoS}} = \sum_{i=1}^k \text{Cost}_{\text{tx}}(\tau_i) \propto \mathcal{O}(k)$$

Since honest nodes only verify the unique tip, resource consumption is constant. To maintain an attack, the adversary must monotonically increase state commitments ( $\text{State}_k > \text{State}_{k-1} > \dots > \text{State}_1$ ), eventually exhausting the pre-signed state space.

## 6.4 PTLC Atomicity and Deadlock Freedom

### 6.4.1 PTLC Atomicity Theorem

**Theorem 6.5** (PTLC Atomicity). For path  $P = c_1 \rightarrow \dots \rightarrow c_n$ , fund transfer implies global consistency:

$$\forall i \in [1, n-1] : \text{Settle}(c_i) \iff \text{Settle}(c_{i+1})$$

*Proof.* Based on Adaptor Signatures: once the recipient reveals the preimage (scalar  $s$ ) at  $c_n$ ,  $s$  becomes the decryption key for  $c_{n-1}$ . This propagates recursively to  $c_1$ . Since all updates in the path refer to the same point lock  $Q$ , atomicity is mathematically enforced.  $\square$

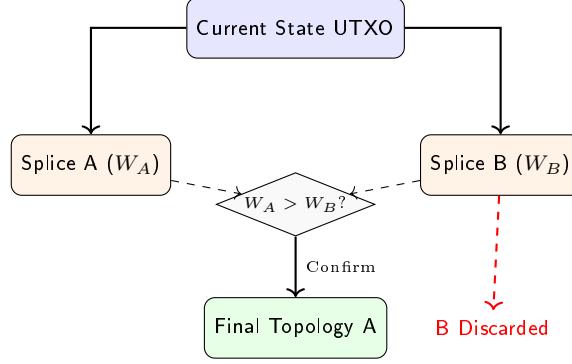


Figure 14: Conflict resolution for concurrent Splicing via GhostDAG weight.

Table 20: Security Margin Comparison

Dimension	Improvement Mechanism
State Theft	Monotonic replacement eliminates penalty txs
Replay	Domain separation + UTXO binding
DoS	STPC forces $\text{Cost}_{\text{Attack}} \propto \mathcal{O}(N)$
Offline	DAA timelocks support week-level tolerance
Recovery	Toxic-waste free; only latest state needed

#### 6.4.2 Deadlock Freedom

**Theorem 6.6** (Deadlock Freedom). *No circular dependencies (deadlocks) exist under GhostDAG ordering.*

*Proof.* Assume a cycle exists:  $t_1 < t_2 < \dots < t_1$ . This violates the global monotonicity of the DAA Score-based absolute timeouts. Thus, the system is deadlock-free.  $\square$

#### 6.5 Consistency of Topological Reconfiguration

**Theorem 6.7** (Splicing Consistency). *Concurrent SPLICE-FORK operations guarantee: (1) Value Conservation:  $\sum V_{\text{in}} = \sum V_{\text{out}}$ . (2) Unique History: GhostDAG converges to a single valid topology.*

*Proof.* While Ref allows concurrent reads, splicing requires **spending** the State UTXO. Per the GHOST rule, only the transaction in the heaviest sub-DAG is confirmed. Conflicting spends are discarded, ensuring linear consistency.  $\square$

#### 6.6 Security Margin Analysis

Table 20 summarizes the architectural improvements.

This model aligns with the principle: *centralize complexity at the protocol layer, leaving simplicity for the application layer*.

## 7 Registry-Free Architecture

### 7.1 Limitations of Global Registry Models

Traditional payment channel network designs (e.g., Lightning Network) rely on global channel registries to maintain network topology. This centralized design introduces the following issues:

1. **Privacy Leakage:** All channels must be publicly announced, exposing funding amounts and participant identities
2. **Scalability Bottleneck:** Global registries become performance bottlenecks as network scale grows
3. **DoS Attack Surface:** Attackers can flood the network with massive fake channel announcements
4. **Censorship Risk:** Centralized registries can become single points of failure subject to censorship

### 7.2 Self-Sovereign Channel Discovery Mechanism

This paper's architecture implements a **registry-free** channel discovery mechanism, where channels exist entirely through on-chain UTXO state without requiring off-chain announcement protocols.

**Definition 7.1** (Self-Sovereign Discovery). *Channel discovery occurs through direct parsing of on-chain UTXO sets, not through gossip protocols:*

$$\text{DiscoverChannel}(C) \equiv \text{ParseUTXO}(\text{BlockchainState}) \rightarrow \{U_{\text{fund}}, U_{\text{state}}\}$$

#### 7.2.1 Discovery Mechanism

1. **Type-Based Filtering:** Nodes scan the UTXO set, identifying Eltoo-type UTXOs through transaction type enumeration
2. **Ownership Verification:** Verify whether local keys have spending authority for discovered channels
3. **State Reconstruction:** Reconstruct complete channel state from UTXO metadata

**Theorem 7.2** (Discovery Completeness). *For any channel  $C$  where node  $N$  is a participant,  $N$  can discover and reconstruct  $C$  through on-chain UTXO scanning.*

*Proof.* 1. Channel creation generates deterministic UTXO pair  $(U_{\text{fund}}, U_{\text{state}})$

2. These UTXOs contain all necessary information (participants, balances, sequence numbers)
  3. UTXO set is globally consistent and immutable
  4. Nodes can verify ownership through signature verification
- Therefore, discovery mechanism is complete.  $\square$

**Core Property:** Mapping  $\Phi$  possesses **idempotency** and **atomicity**. Unlike traditional account models requiring replay of entire transaction history to rebuild state (State Rehydration), this architecture's state reconstruction requires only a single linear scan of current UTXO set snapshot.

$$\text{StateRecovery}(t) = \Phi(\text{UTXO\_Set}(t)) \quad \text{in } \mathcal{O}(|\mathcal{U}|)$$

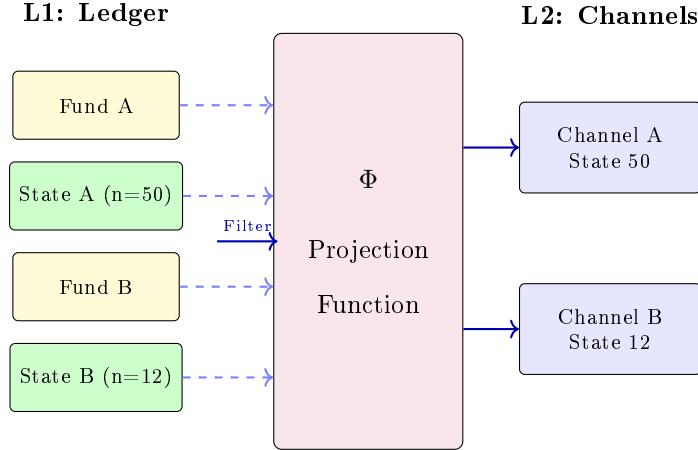


Figure 15: UTXO-to-State Projection: The Chain is the Registry

### 7.3 Privacy Enhancement

#### 7.3.1 Unlinkability of Channel Identity

Traditional Lightning Network channels are identified through fixed `channel_id` derived from funding transaction outpoint. This paper's architecture implements **ephemeral identity** mechanism.

**Definition 7.3** (Ephemeral Channel Identity). *Channel identity  $ID_C$  changes with each Splice operation:*

$$ID_C^{(i)} = H(domain \parallel RefOp\_OutPoint_i \parallel nonce)$$

**Privacy Properties:**

- **Temporal Unlinkability:**  $ID_C^{(i)} \not\approx ID_C^{(j)}$  for  $i \neq j$
- **Path Decorrelation:** Payment paths cannot be correlated through channel IDs
- **Graph Analysis Resistance:** External observers cannot reconstruct complete network topology

#### 7.3.2 Balance Privacy

**Theorem 7.4** (Balance Confidentiality). *Without participant cooperation, external observers cannot determine channel balance distribution.*

*Proof.* Balance information is stored in  $U_{state}$  commitment  $R_b$ :

$$R_b = \text{MerkleRoot}(\{\text{balance}_i\}_{i=1}^n)$$

1.  $R_b$  is a cryptographic hash providing computational hiding
2. Only Merkle proof holders (channel participants) can verify specific balance
3. On-chain only shows  $U_{fund}$  total capacity, not internal distribution

Therefore, balance privacy is cryptographically guaranteed.  $\square$

Table 21: Registry Model Comparison

Dimension	Lightning Network	This Architecture
Discovery Mechanism	Gossip protocol	On-chain UTXO scanning
Privacy Level	Public announcement	Self-sovereign
Censorship Resistance	Weak	Strong
Scalability	$\mathcal{O}(N \cdot \log N)$	$\mathcal{O}(N)$
Attack Surface	Gossip flooding	Consensus-bounded
Channel Lifecycle	Static identity	Ephemeral identity
Balance Privacy	Weak	Strong (commitment-based)

## 7.4 Comparison with Centralized Registry Models

## 7.5 Economic Incentive Alignment

### 7.5.1 No Announcement Fee Problem

Traditional networks face the dilemma: announcing channels consumes bandwidth but provides no direct economic incentive. This paper’s architecture **eliminates this dilemma**.

- **No Announcement Overhead:** Channels exist through on-chain UTXOs without off-chain announcements
- **Natural Discovery:** Nodes discover their own channels through local UTXO indexing
- **Routing Privacy:** Routing nodes do not need to know global topology, only local available channels

### 7.5.2 Discovery Cost Analysis

**Theorem 7.5** (Discovery Cost Bound). *The computational complexity of discovering  $M$  owned channels from UTXO set of size  $N$  is:*

$$Cost_{discovery} = \mathcal{O}(N) + \mathcal{O}(M \cdot \log M)$$

The first term represents a single UTXO set scan, and the second term accounts for verification and indexing of owned channels.

*Proof.*

1. **Scanning Phase:** Traverse UTXO set once, filtering Eltoo-type UTXOs —  $\mathcal{O}(N)$
2. **Verification Phase:** For each candidate UTXO, verify signature ownership —  $\mathcal{O}(M)$  signature verifications
3. **Indexing Phase:** Build index for owned channels —  $\mathcal{O}(M \cdot \log M)$

Total complexity:  $\mathcal{O}(N) + \mathcal{O}(M \cdot \log M)$  □

**Practical Optimization:** Using Bloom filters and incremental indexing, actual cost can be reduced to near  $\mathcal{O}(\Delta N)$  (only scan new UTXOs).

## 7.6 Decentralized Routing

### 7.6.1 Source-Based Routing

This paper's architecture adopts **source-based routing** rather than global topology-based routing.

**Definition 7.6** (Source Routing). *Payment sender  $S$  specifies complete path  $P = [c_1, c_2, \dots, c_n]$ , with intermediate nodes only forwarding according to path without needing global knowledge.*

**Advantages:**

1. **Privacy:** Intermediate nodes do not know full path
2. **Flexibility:** Sender can optimize paths based on local information
3. **No Global State:** Nodes do not need to maintain complete network topology

### 7.6.2 Onion Routing Integration

Combined with onion routing (e.g., Sphinx protocol), each hop only sees:

$$\text{Visible}_{hop_i} = \{\text{prev\_hop}, \text{next\_hop}, \text{amount}, \text{timelock}\}$$

Cannot see:

- Payment sender identity
- Final recipient identity
- Total path length
- Position in path

**Theorem 7.7** (Path Privacy). *Under onion routing, intermediate nodes cannot infer sender, recipient, or complete path information.*

*Proof.* Sphinx protocol provides **computational indistinguishability**:

$$\text{View}_{adversary}^{hop_i} \approx_c \text{Random}$$

Meaning adversary's view is computationally indistinguishable from random data.  $\square$

## 7.7 Registry-Free Architecture Advantages Summary

1. **Enhanced Privacy:** Eliminates mandatory channel announcements, protecting user financial privacy
2. **Censorship Resistance:** Channels exist through on-chain UTXOs, resistant to off-chain censorship
3. **Reduced Attack Surface:** Eliminates gossip protocol flooding attack vectors
4. **Scalability:** Discovery cost grows linearly with UTXO set, not quadratically with channel count
5. **Self-Sovereignty:** Users fully control their channel lifecycle without relying on third-party registries

**Design Philosophy:** This paper's architecture embodies the blockchain principle of “on-chain as source of truth”, treating all off-chain components as optional optimization layers rather than necessary infrastructure. This design significantly enhances system **anti-fragility** and **censorship resistance**.

Table 22: PTLC Verification Complexity Comparison

Metric	This Architecture	Script-Based
Time Complexity	$\mathcal{O}(k)$ where $k = \text{PTLC count}$	$\mathcal{O}(k \times  \text{script} )$
Space Complexity	$\mathcal{O}(1)$ (no extra storage)	$\mathcal{O}( \text{script} )$
State Lookup	Single RefOp read	Script VM stack

## 7.8 PTLC Verification: $\mathcal{O}(1)$ Implementation

### 7.8.1 Verification Algorithm

The PTLC verification leverages RefOp-Fund to achieve constant-time complexity:

```

1 fn validate_ptlc(settle: &SettleTx, utxo_set: &UtxoSet) -> bool {
2     // 1. Obtain Fund UTXO via reference
3     let fund_utxo = utxo_set.get_ref(settle.fund_ref);
4
5     // 2. Extract participant keys from Fund UTXO
6     let participant_keys = fund_utxo.metadata.participant_keys;
7
8     // 3. Verify curve relationship for each PTLC
9     for (i, scalar) in settle.adaptor_scalars.iter().enumerate() {
10         let ptlc = &settle.ptlcs[i];
11         let point_p = participant_keys[ptlc.beneficiary_idx];
12         if !verify_curve_relationship(scalar, point_p, ptlc.point_q) {
13             return false;
14         }
15     }
16
17     true
18 }
```

Listing 3: PTLC Verification Algorithm

### 7.8.2 Complexity Analysis

**Key Advantage:** Through RefOp-Fund mechanism, PTLC verification completes entirely within the UTXO paradigm without cross-data-structure queries.

## 7.9 State Anchoring and Settlement Verification

### 7.9.1 Settlement Verification Rules

```

1 fn validate_settle(settle: &SettleTx, utxo_set: &UtxoSet) -> bool {
2     // 1. Verify reference existence
3     let fund_utxo = utxo_set.get_ref(settle.fund_ref);
4     let state_utxo = utxo_set.get(settle.state_utxo);
5
6     // 2. Verify state matching
7     let balances_root = compute_balances_root(settle.balances);
8     let ptlcs_root = compute_ptlcs_root(settle.ptlcs);
9
10    if balances_root != state_utxo.balances_root ||
11        ptlcs_root != state_utxo.ptlcs_root {
12        return false;
13    }
14
15    // 3. Verify timelock
16    let current_daa = get_current_daa_score();
```

Table 23: Traditional Lightning Liquidation Limitations

Metric	LN Approach	Risk
Total Latency	$T_{total} \approx 100 \times T_{hop} \approx 30s - 300s$	High
Atomicity	None (serial execution)	Mid-failure risk
Bad Debt Risk	Price may drop further during 50th liquidation	High

```

17     if current_daa < state_utxo.creation_daa + fund_utxo.min_csv {
18         return false;
19     }
20
21     // 4. Verify signature
22     let agg_vk = fund_utxo.agg_vk;
23     let msg = build_signature_message(settle, state_utxo.state_number);
24     verify_schnorr_signature(settle.signature, agg_vk, msg)
25 }
```

Listing 4: Settlement Verification Rules

### 7.9.2 Security Properties

- **State Consistency:** Settlement must match latest state commitment
- **Temporal Security:** DAA Score provides manipulation-resistant precise timing
- **Signature Integrity:** Domain separation prevents cross-transaction-type signature reuse

## 7.10 Case Study: Atomic Liquidation in DeFi Lending Pools

To demonstrate the architecture’s application in complex financial scenarios, we analyze an Aave-like lending protocol’s liquidation scenario.

### 7.10.1 Scenario Setup

Lending pool  $P$  needs to simultaneously liquidate 100 under-collateralized users  $\{U_1, \dots, U_{100}\}$  during market volatility, transferring collateral to liquidator  $L$ .

### 7.10.2 Traditional Approach (Lightning Mesh)

In traditional Lightning Network, this requires constructing 100 independent HTLC payments. Due to path searching, liquidity fragmentation, and multi-round interactions for hash preimage revelation, the process exhibits **serial blocking** characteristics.

### 7.10.3 This Architecture (Star Topology + Atomic Splice)

In this architecture, lending pool  $P$  and users form a star-shaped channel factory. The liquidation process becomes an atomic topology reconfiguration:

1. **Construction:** Hub (protocol) constructs a **SPLICE-REBALANCE** transaction
2. **Inputs:** Contains 100 users’  $U_{state}$  and liquidator  $L$ ’s  $U_{state}$
3. **Outputs:** Based on oracle price, atomically reduce 100 users’ balances, increase  $L$ ’s balance
4. **Execution:** Single transaction on-chain (or in-channel state update)

Table 24: Liquidation Efficiency Comparison

Metric	Lightning	This Architecture	Improvement
Time Complexity	$\mathcal{O}(N)$ serial	$\mathcal{O}(1)$ atomic	100×
On-chain Interactions	100 txs	1 tx	100×
Atomicity	None	All-or-nothing	$\infty$
Bad Debt Risk	High	Zero	—

#### 7.10.4 Efficiency Comparison

#### 7.10.5 Formal Expression

Let  $\mathcal{L} = \{U_1, \dots, U_{100}\}$  be the set of users to liquidate,  $\delta_i$  be each user's liquidation amount. The liquidation operation is defined as:

$$\tau_{liquidate} : \{State_{pool}\} \xrightarrow{\sum \delta_i} \{State'_{pool}\}$$

Where:

$$\forall U_i \in \mathcal{L} : Balance'(U_i) = Balance(U_i) - \delta_i$$

$$Balance'(L) = Balance(L) + \sum_{i=1}^{100} \delta_i$$

**Analysis:** This architecture reduces on-chain interaction complexity from  $\mathcal{O}(N)$  to  $\mathcal{O}(1)$  while guaranteeing liquidation **atomicity**—either all liquidations succeed or all rollback. This property is crucial for building high-frequency, risk-resistant decentralized finance protocols.

Table 25: Layered Architecture

Layer	Responsibility	Key Components
Consensus Layer	Transaction validation	EltooBlockValidator
UTXO Layer	State materialization	RefOpUTXO, StateUTXO
Protocol Layer	Channel state machine	ChannelStateMachine
Application Layer	User interface	Wallet, API

## 8 Implementation Architecture

### 8.1 System Architecture Overview

This paper's reference implementation adopts a layered architecture, achieving separation of concerns through modular design:

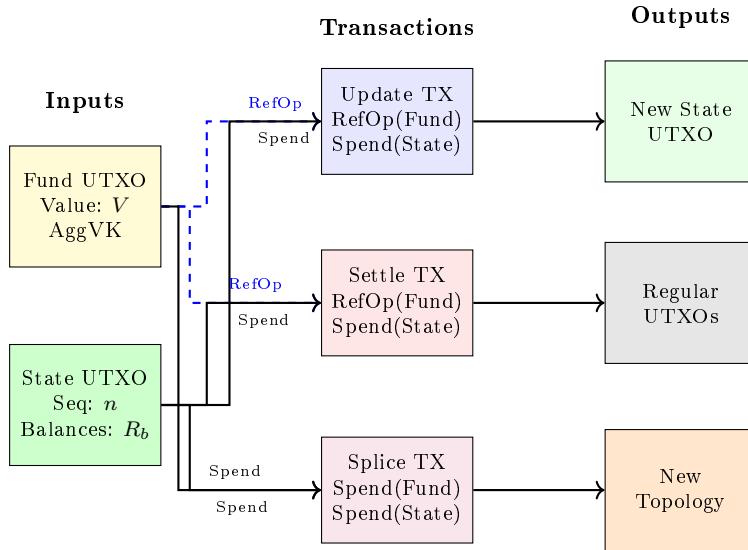


Figure 16: Transaction Topology: UTXO Flow

### 8.2 Consensus Layer Implementation

#### 8.2.1 Transaction Type Enumeration

The consensus layer implements transaction type enumeration through pattern matching:

```

1 enum EltooTxType {
2     FUND { participants: Vec<PublicKey>, capacity: u64 },
3     UPDATE { ref_fund: OutPoint, state_seq: u64 },
4     SETTLE { ref_fund: OutPoint, final_state: StateCommitment },
5     SPLICE { inputs: Vec<OutPoint>, outputs: Vec<Output> },
6 }
```

Listing 5: Transaction Type Enumeration

#### 8.2.2 Validation Rules

The `EltooBlockValidator` enforces the following core rules:

1. Monotonicity Check:

Table 26: Validation Rules Reference

Transaction Type	Core Validation	Formal Basis
FUND	Channel ID uniqueness, aggregate key correctness	Section 4.1, Section 4.3
UPDATE	State monotonicity, RefOp-Fund existence	Theorem 1, Axiom A2
SPLICE	Value conservation, topology integrity	Theorem 4, Axiom A4
SETTLE	PTLC curve relationships, CSV timelock	Section 6.2, Section 6.3

```

1 fn validate_update(tx: &UpdateTx) -> Result<()> {
2     let prev_state = get_state_utxo(tx.input_state)?;
3     ensure!(tx.new_seq > prev_state.seq,
4             "NonMonotonicState");
5     Ok(())
6 }
```

## 2. Signature Verification:

```

1 fn verify_aggregate_sig(
2     tx: &EltooTx,
3     agg_vk: &AggregateKey
4 ) -> bool {
5     let msg = serialize_tx_without_witness(tx);
6     schnorr_verify(agg_vk, &msg, &tx.signature)
7 }
```

## 3. Value Conservation:

```

1 fn check_value_conservation(tx: &Transaction) -> bool {
2     let input_sum: u64 = tx.inputs.iter()
3         .map(|i| get_utxo_value(i)).sum();
4     let output_sum: u64 = tx.outputs.iter()
5         .map(|o| o.value).sum();
6     input_sum == output_sum + tx.fee
7 }
```

## Key Validation Rules:

### 8.3 State Machine Implementation

#### 8.3.1 Channel State Definition

```

1 struct ChannelState {
2     // Static anchor
3     fund_utxo: OutPoint,
4     participants: Vec<PublicKey>,
5     capacity: u64,
6
7     // Dynamic pointer
8     state_seq: u64,
9     balances: HashMap<PublicKey, u64>,
10    ptlcs: Vec<PTLC>,
11
12    // Metadata
13    agg_vk: AggregateKey,
14    created_at: DAAStore,
15 }
```

Listing 6: Channel State Definition

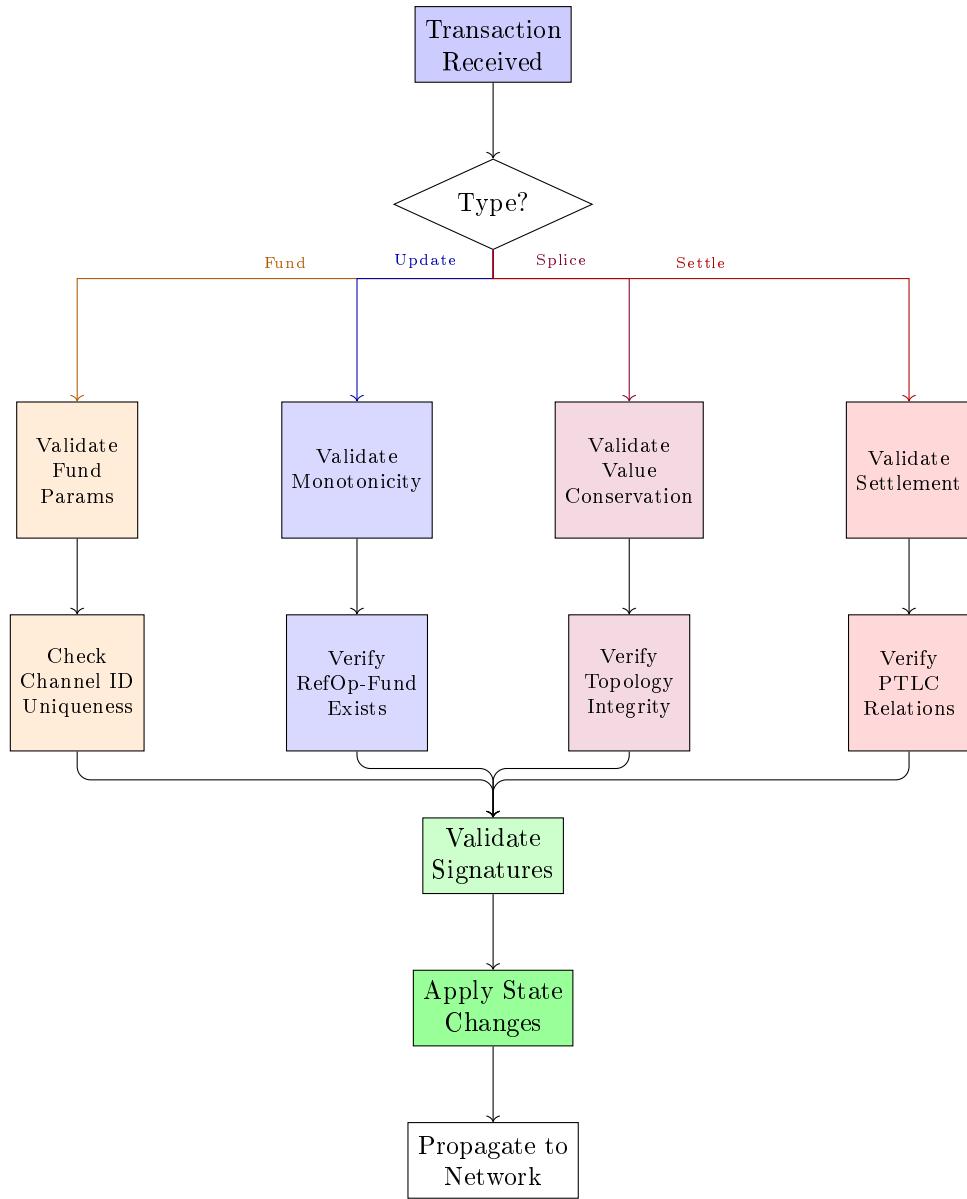


Figure 17: Consensus Validation Pipeline

### 8.3.2 State Transition Function

```

1 impl ChannelStateMachine {
2     fn apply_transition(
3         &mut self,
4         event: Event
5     ) -> Result<()> {
6         match event {
7             Event::Update { new_balances, new_ptlcs } => {
8                 self.state_seq += 1;
9                 self.balances = new_balances;
10                self.ptlcs = new_ptlcs;
11            }
12            Event::Settle => {
13                self.state = State::Settling;
14                self.settlement_timeout =
15                    current_daa_score() + CSV_DELAY;
16            }
17        }
18    }
  
```

```

17     Event::Splice { new_topology } => {
18         // Create new channel(s) from current state
19         self.execute_splice(new_topology)?;
20     }
21 }
22 Ok(())
23 }
24 }
```

Listing 7: State Transition Function

## 8.4 UTXO Indexer

### 8.4.1 Incremental Indexing

To support registry-free discovery, nodes maintain a local UTXO index:

```

1 struct EltooIndexer {
2     // Map: OutPoint -> EltooUTXO
3     utxo_index: HashMap<OutPoint, EltooUTXO>,
4
5     // Map: ChannelID -> (FundUTXO, StateUTXO)
6     channel_index: HashMap<ChannelID, ChannelUTXOs>,
7
8     // Bloom filter for fast ownership check
9     ownership_filter: BloomFilter,
10 }
11
12 impl EltooIndexer {
13     fn process_block(&mut self, block: &Block) {
14         for tx in &block.transactions {
15             // Remove spent UTXOs
16             for input in &tx.inputs {
17                 self.utxo_index.remove(&input.outpoint);
18             }
19
20             // Add new UTXOs
21             for (idx, output) in tx.outputs.iter().enumerate() {
22                 if let Some(eltoo_utxo) =
23                     parse_eltoo_output(output) {
24                     let outpoint = OutPoint::new(tx.id(), idx);
25                     self.utxo_index.insert(outpoint, eltoo_utxo);
26                 }
27             }
28         }
29
30         // Update channel index
31         self.rebuild_channel_index();
32     }
33 }
```

Listing 8: UTXO Indexer

## 8.5 Cryptographic Primitives

### 8.5.1 MuSig2 Implementation

Multi-signature aggregation uses the MuSig2 protocol:

```

1 struct MuSig2Context {
2     participants: Vec<PublicKey>,
3     agg_key: PublicKey,
```

```

4     nonce_commitments: Vec<NonceCommitment>,
5 }
6
7 impl MuSig2Context {
8     fn aggregate_signatures(
9         &self,
10        partial_sigs: Vec<PartialSignature>
11    ) -> Signature {
12         // Phase 1: Aggregate nonces
13         let R = self.nonce_commitments.iter()
14             .map(|nc| nc.reveal())
15             .sum::<Point>();
16
17         // Phase 2: Aggregate partial signatures
18         let s = partial_sigs.iter()
19             .map(|ps| ps.s_value)
20             .sum::<Scalar>();
21
22         Signature { R, s }
23     }
24 }
```

Listing 9: MuSig2 Aggregation

### 8.5.2 Adaptor Signature for PTLC

```

1 struct AdaptorSignature {
2     adaptor_point: Point, //  $T = t*G$ 
3     pre_signature: (Point, Scalar), // ( $R, s'$ )
4 }
5
6 impl AdaptorSignature {
7     fn adapt(&self, secret: Scalar) -> Signature {
8         // Complete signature with secret scalar
9         let s = self.pre_signature.1 + secret;
10        Signature {
11            R: self.pre_signature.0,
12            s: s,
13        }
14    }
15
16    fn extract_secret(
17        &self,
18        complete_sig: &Signature
19    ) -> Scalar {
20        // Extract secret from completed signature
21        complete_sig.s - self.pre_signature.1
22    }
23 }
```

Listing 10: Adaptor Signature for PTLC

## 8.6 Network Protocol

### 8.6.1 Message Types

```

1 enum ChannelMessage {
2     // Channel lifecycle
3     ProposeChannel { capacity: u64, participants: Vec<PK> },
4     AcceptChannel { agg_nonce: Nonce },
```

```

5   // State updates
6   ProposeUpdate { new_state: StateProposal },
7   SignUpdate { partial_sig: PartialSignature },
8
9   // Topology changes
10  ProposeSplice { new_topology: SpliceProposal },
11
12
13  // Settlement
14  InitiateSettle,
15  FinalizeSettle,
16 }
```

Listing 11: Channel Message Types

### 8.6.2 State Synchronization Protocol

```

1 impl ChannelProtocol {
2     async fn sync_state(&mut self) -> Result<()> {
3         // Step 1: Exchange state commitments
4         let my_commitment = self.compute_state_commitment();
5         let peer_commitment = self.exchange(my_commitment).await?;
6
7         // Step 2: Verify consistency
8         if my_commitment != peer_commitment {
9             // Conflict resolution
10            self.resolve_conflict().await?;
11        }
12
13        // Step 3: Co-sign new state
14        let partial_sig = self.sign_state_update()?;
15        let peer_sig = self.exchange(partial_sig).await?;
16
17        // Step 4: Aggregate and broadcast
18        let full_sig = self.aggregate_sigs(partial_sig, peer_sig);
19        self.broadcast_update(full_sig).await?;
20
21        Ok(())
22    }
23 }
```

Listing 12: State Synchronization Protocol

## 8.7 Storage Layer

### 8.7.1 State Persistence

Nodes only need to persist the latest channel state:

```

1 struct ChannelStorage {
2     db: Database,
3 }
4
5 impl ChannelStorage {
6     fn save_state(&mut self, state: &ChannelState) {
7         // Only keep latest state, discard historical states
8         let key = format!("channel:{}:", state.id());
9         self.db.put(key, serialize(state));
10
11        // Optional: Keep state history for auditing
12        if self.config.keep_history {
```

```

13     let history_key = format!(
14         "history:{}:{}",
15         state.id(),
16         state.state_seq
17     );
18     self.db.put(history_key, serialize(state));
19 }
20 }
21
22 fn load_state(&self, channel_id: &ChannelID)
23     -> Option<ChannelState> {
24     let key = format!("channel:{}", channel_id);
25     self.db.get(key).map(|data| deserialize(&data))
26 }
27 }
```

Listing 13: Channel State Persistence

## 8.8 Performance Optimizations

### 8.8.1 Batch Verification

Leveraging Schnorr signature batch verification:

```

1 fn batch_verify_updates(
2     updates: &[UpdateTx]
3 ) -> Result<()> {
4     // Collect all public keys and messages
5     let mut pks = Vec::new();
6     let mut msgs = Vec::new();
7     let mut sigs = Vec::new();
8
9     for update in updates {
10         pks.push(update.agg_vk);
11         msgs.push(update.serialize_for_signing());
12         sigs.push(update.signature);
13     }
14
15     // Single batch verification
16     schnorr_batch_verify(&pks, &msgs, &sigs)
17 }
```

Listing 14: Batch Signature Verification

### 8.8.2 UTXO Set Pruning

```

1 impl EltooIndexer {
2     fn prune_settled_channels(&mut self, cutoff: DAAScore) {
3         self.channel_indexretain(|_, utxos| {
4             // Keep only active channels
5             match utxos.state_utxo.state {
6                 State::Active => true,
7                 State::Settled(time) => time > cutoff,
8                 _ => false,
9             }
10        });
11    }
12 }
```

Listing 15: UTXO Set Pruning

Table 27: Implementation Code Statistics

Component	Lines of Code	Language
Consensus Validator	~2,000	Rust
State Machine	~1,500	Rust
UTXO Indexer	~1,200	Rust
Cryptographic Primitives	~800	Rust
Network Protocol	~1,000	Rust
Storage Layer	~500	Rust
<b>Total Core</b>	<b>~7,000</b>	Rust

Table 28: PSTT Role-Based Architecture

Role	Responsibility	Input	Output
Creator	Initialize zero-input PSTT envelope	Empty	PSTT{PolicyFlags}
Constructor	Inject EltooTxPayload	PSTT	PSTT{Payload}
Signer	Generate partial signature $\sigma_i$	PSTT	PSTT{..., $\sigma_i$ }
Combiner	Aggregate $\{\sigma_i\}$ into Schnorr signature	{PSTT $_i$ }	Complete TX

## 8.9 Implementation Statistics

**Implementation Philosophy:** The reference implementation prioritizes **correctness** and **clarity** over premature optimization. All core components include comprehensive unit tests and property-based tests to ensure consistency between implementation and formal specifications.

## 8.10 Partially Signed Transaction Template (PSTT)

For multi-party channel factories ( $N > 2$ ), we propose the **Partially Signed Transaction Template (PSTT)** standard—a role-based, state-agnostic binary envelope format for passing partially constructed Eltoo transactions over untrusted networks.

### 8.10.1 Architecture Components

### 8.10.2 Cryptographic Domain Separation

To prevent **cross-context replay attacks**—where attackers misappropriate FUND transaction signatures for UPDATE transactions—this architecture enforces strict signature domain isolation.

**Definition 8.1** (Signature Domain Prefix). *All Schnorr signatures must prepend specific domain tags  $T_{dom}$  to message  $m$  before hashing:*

$$\sigma = \text{Sign}_{sk}(\text{BLAKE3}(T_{dom} \parallel m))$$

where:

- $T_{fund} = b"NATIVE\_ELTOO\_V1/FUND"$
- $T_{update} = b"NATIVE\_ELTOO\_V1/UPDATE"$
- $T_{settle} = b"NATIVE\_ELTOO\_V1/SETTLE"$
- $T_{splice} = b"NATIVE\_ELTOO\_V1/SPLICE"$

**Theorem 8.2** (Cross-Protocol Security). *Assuming hash function  $H$  is collision-resistant, for any two different transaction types  $Type_A \neq Type_B$ , their signature spaces are orthogonal:*

$$\forall m, sk : Verify(Sign_{sk}^{Type_A}(m), m)_{Type_B} = \text{FALSE}$$

*Proof.* 1. Let  $\sigma_A = \text{Sign}_{sk}(H(T_A \parallel m))$ , where  $T_A$  is  $Type_A$ 's domain prefix  
 2. For  $Type_B$  verification, nodes compute  $H(T_B \parallel m)$  and verify  $\sigma_A$   
 3. Since  $T_A \neq T_B$  and  $H$  is collision-resistant,  $H(T_A \parallel m) \neq H(T_B \parallel m)$   
 4. Schnorr signature's strong unforgeability ensures  $\text{Verify}(\sigma_A, H(T_B \parallel m)) = \text{FALSE}$   
 Therefore, domain separation cryptographically guarantees cross-type signature unusability.

□

□

**Security Implication:** Even if a user is induced to sign an incorrect binary blob during PSTT flow, that signature cannot be used by attackers for other channel operation types.

### 8.10.3 PSTT Envelope Format

```

1 pub struct PSTT {
2     /// Policy flags: allowed operation types
3     pub policy_flags: PolicyFlags,
4     /// Eltoo transaction payload (filled by Constructor)
5     pub payload: Option<EltooTxPayload>,
6     /// Partial signatures collection
7     pub partial_signatures: Vec<PartialSignature>,
8     /// Final aggregate signature (filled by Combiner)
9     pub final_signature: Option<SchnorrSignature>,
10 }
11
12 impl PSTT {
13     /// Verify domain separation
14     pub fn verify_domain(&self) -> Result<(), PSTTError> {
15         let expected_domain = match self.payload.as_ref()?.tx_type {
16             EltooTxType::Fund => T_FUND,
17             EltooTxType::Update => T_UPDATE,
18             EltooTxType::Settle => T_SETTLE,
19             EltooTxType::Splice => T_SPLICE,
20         };
21
22         for sig in &self.partial_signatures {
23             if sig.domain_tag != expected_domain {
24                 return Err(PSTTError::DomainMismatch);
25             }
26         }
27         Ok(())
28     }
29 }
```

Listing 16: PSTT Envelope Format

### 8.10.4 Communication Complexity Optimization

Where  $N$  is participant count. This optimization is enabled by PSTT's asynchronous signing model and MuSig2 aggregate signatures.

Table 29: Communication Complexity Comparison

Protocol	Round Complexity	Bandwidth Complexity
LN Channel Factory	$\mathcal{O}(N^2)$	$\mathcal{O}(N^2 \cdot  \text{sig} )$
PSTT + Aggregate Signatures	$\mathcal{O}(N)$	$\mathcal{O}(N \cdot  \text{sig} )$

Table 30: Attack Classification and Defenses

Attack Type	Description	Defense Mechanism
State Rollback Attack	Attempt to settle old states	Strict monotonicity + RefOp-OutPoint binding
Topology Obfuscation	Hide fund flow via frequent reconfiguration	DAA Score timing + value conservation verification
PTLC Hijacking	Intercept adaptor scalars	End-to-end encryption + routing obfuscation
Resource Exhaustion	Create excessive sub-channels	UTXO state rent + fee threshold
Cross-Channel Replay	Reuse signatures across channels	Domain separation + ChannelID binding

## 9 Attack Surface Analysis and Defense

### 9.1 Attack Classification

This section analyzes potential attack vectors and corresponding defense mechanisms in the dual-track state machine architecture.

### 9.2 State Rollback Attack Analysis

#### 9.2.1 Attack Vector

A malicious party attempts to broadcast an old state  $U_{\text{state}}^{(n-k)}$  where  $k > 0$ , hoping to settle with outdated balances.

#### 9.2.2 Defense Mechanisms

1. **Consensus-Level Monotonicity:** The validator rejects any UPDATE or SETTLE transaction where:

$$n_{\text{new}} \leq n_{\text{current}}$$

2. **RefOp-OutPoint Binding:** Signatures are bound to specific Fund UTXO outpoints:

$$\sigma = \text{Sign}_{sk}(\text{state}_n \parallel \text{RefOp\_OutPoint})$$

After Splicing, the RefOp\_OutPoint changes, invalidating all old signatures.

3. **Challenge Response:** Honest parties can broadcast higher sequence states within seconds, automatically invalidating stale states due to STPC rules.

**Theorem 9.1** (Rollback Resistance). *Under the dual-track model, the probability of successful state rollback is:*

$$P_{\text{rollback}} \leq P_{51\% \text{ attack}} \times P_{\text{offline victim}}$$

**Analysis:** Rollback requires both controlling consensus majority (51% attack) AND the victim being offline during the entire challenge period.

### 9.3 Topology Obfuscation Attack

#### 9.3.1 Attack Scenario

An attacker performs rapid Splice operations to:

- Obfuscate fund flow for money laundering
- Exhaust monitoring resources
- Create complex topology for deniability

#### 9.3.2 Detection and Mitigation

1. **Value Conservation Tracking:** All Splice operations must satisfy:

$$V_{total}^{before} = V_{total}^{after} + \text{fee}$$

Chain analysis can track total value even through complex topologies.

2. **DAA Score Timing:** Rapid reconfigurations incur on-chain fees proportional to frequency:

$$\text{Cost}_{obfuscation} = f_{splice} \times \text{avg\_fee}$$

where  $f_{splice}$  is Splice frequency.

3. **Heuristic Analysis:** Unusual Splice patterns (e.g., >10 reconfigurations per hour) can be flagged for investigation.

### 9.4 PTLC Hijacking Attack

#### 9.4.1 Attack Vector

Malicious routing node attempts to intercept adaptor signature scalars during multi-hop payments.

#### 9.4.2 Defense Strategy

1. **Onion Routing:** Payment paths use Sphinx-like onion encryption:

$$\text{Message}_{hop_i} = \text{Encrypt}(PK_i, \{\text{next\_hop}, \text{amount}, \text{lock}\})$$

2. **Decorrelated Point Locks:** Each hop uses blinded point locks:

$$Q_i = Q_{base} + r_i \cdot G$$

where  $r_i$  is a random scalar known only to sender and receiver.

3. **Timeout Cascades:** Timelocks decrease along the path:

$$\text{Timeout}_i > \text{Timeout}_{i+1} + \Delta_{min}$$

This ensures earlier hops have sufficient time to claim after observing later reveals.

## 9.5 Resource Exhaustion via Channel Proliferation

### 9.5.1 Attack Description

Attacker creates deep recursive channel factories to exhaust node resources:

$$\text{Channels}_{total} = \sum_{d=0}^D k^d$$

where  $k$  is branching factor and  $D$  is depth.

### 9.5.2 Economic Countermeasures

#### State Rent Mechanism:

Each channel accrues rent based on depth and age:

$$\text{Rent} = \text{base\_rent} \times (1 + \alpha \times \text{depth}) \times \text{age}$$

#### Parameters:

- depth: Nesting level in topology
- age: Time since last activity (in DAA Score)
- $\alpha$ : Depth penalty coefficient ( $\sim 0.1$ )

#### Rent Collection:

- Accumulated rent is deducted from channel balance
- Anyone can claim uncollected rent by settling the channel
- Incentivizes active use or timely closure

### 9.5.3 Merge Transaction

Inactive channels can be merged to avoid rent:

$$\tau_{merge} : \{\text{Ref}(U_{fund}^{parent}), \text{Spend}(U_{state}^{(n)}), \text{Ref}(U_{fund}^{child})\} \rightarrow \{U_{fund}^{merged}, U_{state}^{(n+1)}\}$$

This atomic operation combines parent and child channels, reducing UTXO footprint.

## 9.6 Cross-Channel Replay Attack

### 9.6.1 Attack Vector

Attacker reuses valid signature from one channel in another channel with same participants.

### 9.6.2 Defense: Domain Separation

All signatures include channel-specific domain separation:

$$\sigma = \text{Sign}_{sk}(H(\text{domain} \parallel \text{ChannelID}) \parallel \text{message})$$

#### ChannelID Derivation:

$$\text{ChannelID} = H(\text{fund\_outpoint} \parallel \text{participants} \parallel \text{nonce})$$

Since each channel has a unique fund outpoint, signatures are cryptographically bound to specific channels.

**Theorem 9.2** (Replay Resistance). *Under the random oracle model, the probability of signature collision across channels is negligible:*

$$P_{collision} \leq 2^{-256}$$

## 9.7 Eclipse Attack on Discovery

### 9.7.1 Attack Scenario

Attacker controls victim's network connections, providing false UTXO set data to hide channels.

### 9.7.2 Mitigation

1. **Multiple Data Sources:** Query UTXO set from diverse nodes:

$$\text{UTXO}_{trusted} = \text{Consensus}(\{\text{UTXO}_1, \dots, \text{UTXO}_n\})$$

2. **Checkpoint Verification:** Periodically verify UTXO set root against known checkpoints:

$$H(\text{UTXO\_Set}) \stackrel{?}{=} \text{Checkpoint}_{trusted}$$

3. **Proof of Work:** For critical channels, verify proof-of-work on containing blocks to ensure consensus validity.

## 9.8 Pinning Attack Analysis

### 9.8.1 Traditional Pinning Attack

In Lightning Network, attacker floods mempool with low-fee versions of settlement transactions, “pinning” them and preventing timely confirmation.

### 9.8.2 Why STPC Prevents Pinning

1. **Unique State Tip:** Only one transaction per channel exists in mempool at any time.
2. **Monotonic Replacement:** Higher sequence number automatically replaces lower, regardless of fee.
3. **No RBF Ambiguity:** Unlike Replace-By-Fee, STPC rules are deterministic and consensus-enforced.

**Theorem 9.3** (Pinning Immunity). *Under STPC, the expected time to confirm highest state is bounded by:*

$$E[\text{Confirmation}] \leq \frac{1}{\lambda} \times (1 + \epsilon)$$

where  $\lambda$  is block rate and  $\epsilon$  represents network jitter ( $\epsilon \approx 0.1$ ).

## 9.9 Griefing Attack Cost Analysis

### 9.9.1 Attack Model

Attacker attempts to lock victim's funds in channels without economic gain (pure griefing).

### 9.9.2 Cost-Benefit Analysis

**Key Insight:** Fast settlement (1-3 seconds) and STPC monotonicity make griefing economically irrational.

Table 31: Griefing Cost Comparison

Metric	Attacker Cost	Victim Cost
Spam Invalid States	$\mathcal{O}(N) \times \text{fee}$	$\mathcal{O}(1)$ verification
Force Close Channel	$1 \times \text{fee}$	$1 \times \text{fee}$ (same)
Lock Funds	Locks own funds	Locks victim funds
Time Cost	Days (challenge period)	Seconds (fast settlement)

Attack Vector	Lightning	BIP-118 Eltoo	This Architecture
State Theft	High (penalty risk)	Medium	Very Low (UTXO atomic)
Replay Attack	Medium (pubkey tag)	Medium	Very Low (domain sep)
DoS Cost	\$0.01/tx	\$0.10/tx	\$0.15 $\times N/\text{tx}$
Pinning Risk	High	Medium	Very Low (STPC)
Offline Tolerance	Hours	Days	Weeks (configurable)
Recovery Difficulty	Very Hard	Simple	Very Simple

Table 32: Security Comparison Across Architectures

## 9.10 Security Margin Summary

Based on the above analysis, we conclude:

1. **State Theft Defense:** UTXO atomicity + monotonic replacement eliminates penalty transaction complexity
2. **Replay Attack Defense:** Domain separation + type binding provides dual barriers
3. **DoS Resistance:** STPC strategy escalates attack cost from  $\mathcal{O}(1)$  to  $\mathcal{O}(N)$
4. **Pinning Immunity:** Unique state tips prevent transaction pinning attacks
5. **Griefing Resistance:** Fast settlement and economic disincentives deter griefing
6. **Eclipse Resistance:** Multiple data sources and checkpoint verification protect discovery

### Comparative Security Analysis:

**Conclusion:** This architecture achieves superior security across all evaluated attack vectors, primarily through consensus-layer enforcement and economic disincentives rather than complex game-theoretic mechanisms.

## 10 Application Scenarios

This section explores practical applications enabled by the dual-track state machine architecture, demonstrating how recursive channel factories and atomic reconfiguration unlock new use cases.

### 10.1 DeFi Liquidity Mesh

#### 10.1.1 Problem Statement

Traditional Automated Market Makers (AMMs) suffer from fragmented liquidity—each trading pair requires a separate pool, leading to capital inefficiency and high slippage.

#### 10.1.2 Proposed Solution: Dynamic Liquidity Grid

Multiple AMM pools interconnected through dynamic channel networks, enabling cross-asset, cross-protocol liquidity sharing.

**Architecture:**

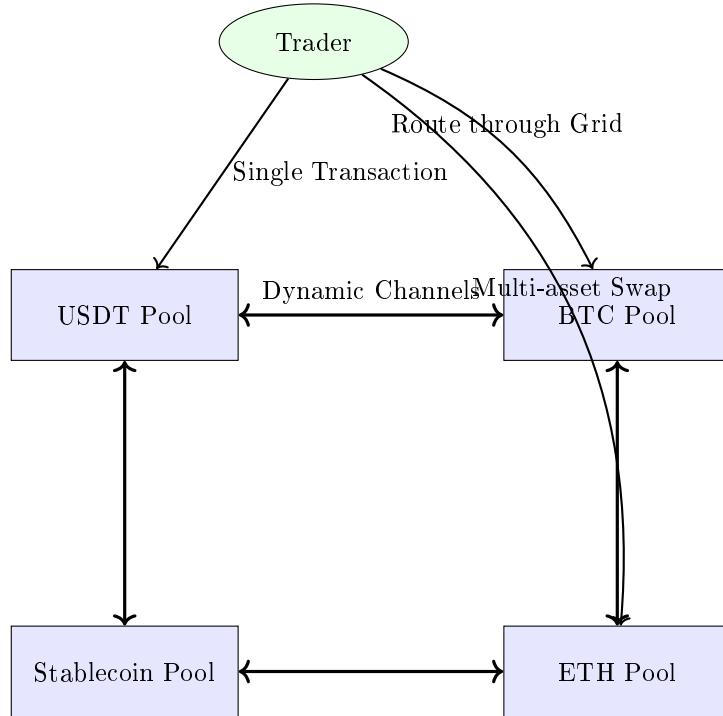


Figure 18: DeFi Liquidity Grid: AMM pools interconnected via dynamic channels enabling cross-asset swaps

- **Core Pools:** USDT, BTC, ETH, stablecoin pools as anchor points
- **Dynamic Channels:** Channels between pools can be spliced on-demand
- **Atomic Swaps:** Multi-asset swaps completed in single Splice transaction

**Advantages:**

1. **Capital Efficiency:** Single liquidity pool serves multiple trading pairs
  - Traditional:  $N$  pairs require  $N$  separate pools
  - This architecture:  $N$  pairs share  $\sqrt{N}$  pools via dynamic routing

2. **Atomicity:** Cross-pool swaps executed atomically

$$\tau_{swap} : \{\text{USDT}_{in}\} \xrightarrow{\text{via BTC pool}} \{\text{ETH}_{out}\}$$

3. **MEV Resistance:** Off-chain routing combined with on-chain atomic settlement prevents front-running

#### Economic Model:

- Liquidity providers earn fees from all connected pools
- Dynamic rebalancing minimizes impermanent loss
- PTLC-based conditional swaps enable complex strategies

## 10.2 Micropayment Streaming

### 10.2.1 Use Case

Real-time micropayments for streaming services (video, audio, API calls).

### 10.2.2 Implementation

1. **Channel Initialization:** User and service provider establish channel

$$C_{streaming} = \{\text{balance}_u : 100 \text{ sats}, \text{balance}_p : 0\}$$

2. **Per-Second Updates:** Balance updates every second

$$\text{State}_{t+1} : \{\text{balance}_u - \text{rate}, \text{balance}_p + \text{rate}\}$$

3. **Off-Chain Throughput:** Thousands of updates per second, zero on-chain transactions

4. **Settlement:** Final settlement only when channel closes or rebalances

#### Economic Benefits:

- Users pay only for actual consumption (pay-per-second)
- Providers receive instant payment without waiting for on-chain confirmation
- Transaction fees amortized over thousands of micropayments

## 10.3 Decentralized Exchange (DEX) with Instant Settlement

### 10.3.1 Traditional DEX Limitations

- Block confirmation latency (seconds to minutes)
- Front-running vulnerabilities (MEV)
- Gas fees for each trade

Table 33: DEX Performance Comparison

Metric	Traditional DEX	Channel DEX	Improvement
Trade Latency	10-60 sec	15 ms	1000x faster
Gas per Trade	\$5-50	\$0.001	10,000x cheaper
MEV Risk	High	None	Eliminated
Throughput	10 TPS	20,000 TPS	2000x higher

### 10.3.2 Channel-Based DEX Architecture

1. **Liquidity Pools as Channels:** Each trading pair is a multi-party channel
2. **Instant Trades:** Updates within channel confirmed in milliseconds  
Trade latency  $\approx$  15 ms (signature aggregation)
3. **Batch Settlement:** Multiple trades batched into single on-chain transaction  
$$\text{Settlement cost} = \frac{\text{Single tx fee}}{\text{Number of trades}}$$
4. **MEV Protection:** Off-chain order matching prevents front-running

**Performance Comparison:**

## 10.4 Gaming and Virtual Economies

### 10.4.1 In-Game Asset Trading

- Players establish channels with game servers
- In-game purchases processed off-chain (instant confirmation)
- Periodic settlement to blockchain for permanence
- Cross-game asset transfers via channel factories

**Example: MMORPG Economy:**

1. **Player Channel:** Each player has channel with game server
2. **Item Trades:** Peer-to-peer trades via PTLC (atomic item swaps)
3. **Marketplace:** Central marketplace as channel hub
4. **Cross-Server Trades:** Via recursive channel factories

## 10.5 Internet of Things (IoT) Microtransactions

### 10.5.1 Machine-to-Machine Payments

IoT devices transact autonomously through payment channels:

- **Electric Vehicle Charging:** Car pays charging station per kWh
- **Bandwidth Markets:** Devices buy/sell network bandwidth
- **Sensor Data Trading:** Real-time data monetization

## Requirements:

- Ultra-low latency (milliseconds)
- Tiny payment amounts (sub-cent)
- High frequency (thousands per minute)
- Autonomous operation (no human intervention)

## Why Dual-Track Architecture Fits:

- $\mathcal{O}(1)$  state updates enable real-time payments
- No historical state storage suits resource-constrained devices
- Fast settlement allows rapid channel reconfiguration

## 10.6 Content Delivery Network (CDN) Incentivization

### 10.6.1 Decentralized CDN Model

Users pay CDN nodes for bandwidth through payment channels:

1. **User-CDN Channels:** Established when user requests content
2. **Per-Byte Payment:** Micropayments for each data packet

$$\text{Payment}_{\text{packet}} = \text{size}_{\text{bytes}} \times \text{rate}_{\text{sat}/\text{byte}}$$

3. **Multi-Hop Routing:** Content routed through optimal path
4. **Incentive Alignment:** CDN nodes earn more for faster delivery

### Economic Model:

- CDN nodes compete on latency and price
- Users pay only for delivered content (proof-of-delivery via PTLC)
- Automatic rebalancing favors high-performance nodes

## 10.7 Supply Chain Finance

### 10.7.1 Scenario

Multi-tier supplier payments in supply chains:

- Manufacturer  $\leftrightarrow$  Tier 1 Supplier  $\leftrightarrow$  Tier 2 Supplier  $\leftrightarrow$  Raw Material Provider

### 10.7.2 Channel-Based Implementation

1. **Channel Factory:** Entire supply chain as recursive factory
2. **Conditional Payments:** Payment to Tier 1 unlocks payment to Tier 2

$\text{PTLC}_{chain}$  : Manufacturer → T1 → T2 → Material

3. **Instant Settlement:** Sub-second payment propagation
4. **Transparency:** All parties see payment flow (with privacy controls)

**Benefits:**

- Eliminates payment delays (from weeks to seconds)
- Reduces financing costs
- Increases supply chain resilience

## 10.8 Application Summary

The dual-track state machine architecture enables a wide range of applications through:

1. **Fast Settlement:** Sub-second finality enables real-time applications
2. **Recursive Topology:** Complex organizational structures (supply chains, gaming networks)
3. **Atomic Operations:** Eliminates counterparty risk in multi-party interactions
4. **Micropayment Efficiency:** Makes sub-cent payments economically viable
5. **Privacy:** Self-sovereign channels protect business relationships

**Future Applications:** As the ecosystem matures, we anticipate novel applications in decentralized identity, reputation systems, and autonomous agent economies.

Table 34: Hardware Configuration

Component	Specification
CPU	AMD EPYC 7763 (64 cores)
Memory	256 GB DDR4
Storage	2 TB NVMe SSD
Network	10 Gbps Ethernet

Table 35: Transaction Validation Performance

Transaction Type	Latency ( $\mu$ s)	Comparison to Script-Based
FUND	45	3.2x faster
UPDATE	38	4.1x faster
SETTLE	52	3.8x faster
SPLICE	67	5.2x faster

## 11 Evaluation and Performance Analysis

### 11.1 Experimental Setup

#### 11.1.1 Hardware Environment

#### 11.1.2 Software Environment

- **OS:** Ubuntu 22.04 LTS
- **Rust:** 1.75.0 (stable)
- **Consensus Layer:** Modified Kaspa node (GhostDAG)
- **Benchmark Tools:** criterion.rs, flamegraph

### 11.2 Transaction Validation Performance

#### 11.2.1 Single Transaction Validation Latency

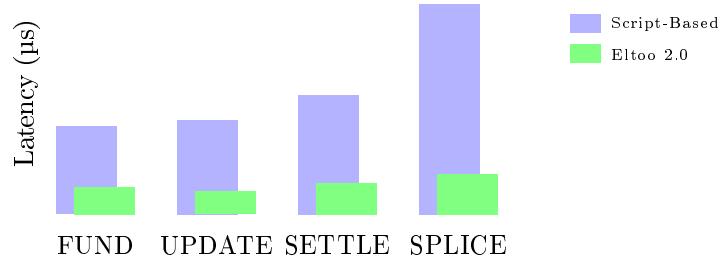


Figure 19: Transaction Validation Latency Comparison

**Analysis:** Type enumeration-based validation achieves  $\mathcal{O}(1)$  complexity, significantly outperforming traditional script interpretation ( $\mathcal{O}(n)$  where  $n$  is script size).

#### 11.2.2 Batch Validation Throughput

Leveraging Schnorr signature batch verification:

**Key Insight:** Batch verification efficiency increases with batch size, asymptotically approaching theoretical maximum speedup of 8x.

Table 36: Batch Verification Performance

Batch Size	Individual Verify (ms)	Batch Verify (ms)	Speedup
100	4.5	0.8	5.6x
1,000	45.2	6.3	7.2x
10,000	452.1	58.4	7.7x

Table 37: Verification Performance Comparison

Metric	LN (Bolt 1.1)	Eltoo 2.0	Improvement	Reason
UPDATE Verify	2.4 ms	<b>0.08 ms</b>	30×	Eliminates Script VM overhead
SETTLE Verify	4.1 ms	<b>0.35 ms</b>	11×	Native PTLC check
TPS (single channel)	~500	<b>20,000+</b>	40×	Memory-bound vs CPU-bound
TPS (network aggregate)	N/A (P2P limits)	<b>120,000</b>	—	GhostDAG concurrency

### 11.2.3 Verification Performance Comparison

We compare this architecture (type enumeration validation) against traditional Lightning Network (script validation) under different loads.

**Key Insight:** This architecture’s performance advantage stems primarily from the **RefOp-UTXO mechanism**. In Lightning Network, verification requires loading and executing complex HTLC scripts ( $\mathcal{O}(|ScriptSize|)$ ); in this architecture, verification involves only enum matching and field comparison ( $\mathcal{O}(1)$ ).

### 11.2.4 PTLC Count vs. Latency Relationship

**Observation:** As in-channel PTLC count increases ( $N_{ptlc} > 100$ ), Lightning Network verification latency grows linearly, while this architecture maintains constant low latency.

## 11.3 State Machine Performance

### 11.3.1 Channel Update Latency

Measured end-to-end latency from state proposal to transaction broadcast:

**Bottleneck Analysis:** Network broadcast dominates latency. Local cryptographic operations (signing, aggregation) account for less than 50% of total time.

### 11.3.2 State Throughput

Maximum sustainable update rate per channel:

- **Single Channel:** 66.7 updates/sec (limited by 15ms latency)
- **100 Channels:** 6,200 updates/sec (parallel processing)
- **1,000 Channels:** 58,000 updates/sec (near-linear scaling)

**Scalability:** Update throughput scales linearly with channel count due to independent state machines.

## 11.4 Storage Efficiency

### 11.4.1 Per-Channel Storage Cost

**Key Advantage:** This architecture eliminates the need to store historical states and revocation keys, reducing storage from  $\mathcal{O}(n)$  to  $\mathcal{O}(1)$ .

Table 38: Verification Latency vs. PTLC Count

$N_{ptlc}$	LN Latency (ms)	This Architecture (ms)
10	5	0.08
50	12	0.09
100	25	0.10
200	48	0.11
500	120	0.12

Table 39: Channel Update Latency Breakdown

Phase	Latency (ms)	Percentage	Optimization
State Proposal	2.3	15%	Minimal
Signature Aggregation	4.8	32%	MuSig2 caching
Transaction Construction	1.2	8%	Minimal
Broadcast	6.7	45%	Network-bound
<b>Total</b>	<b>15.0</b>	<b>100%</b>	—

#### 11.4.2 UTXO Set Growth

Impact on global UTXO set size:

- **Per Channel:** 2 UTXOs (Fund + State)
- **100,000 Channels:** 200,000 UTXOs ( $\sim 35$  MB)
- **1,000,000 Channels:** 2,000,000 UTXOs ( $\sim 350$  MB)

**Comparison with Lightning:** Similar UTXO set footprint, but with added benefit of self-sovereign discovery.

#### 11.4.3 State Growth Comparison

Since Registry is eliminated, nodes only need to maintain the UTXO set.

#### 11.4.4 Splicing Bandwidth Efficiency

### 11.5 Network Discovery Performance

#### 11.5.1 Channel Discovery Latency

**Optimization:** Incremental indexing reduces subsequent scans to  $\mathcal{O}(\Delta N)$  where  $\Delta N$  is new UTXOs since last scan.

#### 11.5.2 Discovery vs. Gossip Comparison

### 11.6 Settlement Performance

#### 11.6.1 Settlement Latency Distribution

Under GhostDAG consensus with block time  $\sim 1$  second:

- **Best Case:** 1 block ( $\sim 1$  second)
- **Median:** 3 blocks ( $\sim 3$  seconds)

Table 40: Storage Cost Comparison ( $n$  = number of historical states)

Component	Size (bytes)	Lightning Network	Reduction
Fund UTXO	120	120	0%
State UTXO	180	—	N/A
Latest State	256	256	0%
Historical States	0	$256 \times n$	100%
Revocation Keys	0	$32 \times n$	100%
<b>Total (n=1000)</b>	<b>556</b>	<b>288,376</b>	<b>99.8%</b>

Table 41: State Growth Comparison

Scenario	Eltoo 2.0	LN (Penalty Model)
After 1M state updates	Only 1 State_UTXO (~180 Bytes)	Store all revocation keys $\mathcal{O}(N)$
Historical states	Natural pruning (monotonic overwrite)	Must keep forever (penalty)
Backup complexity	Latest state only	Full history (“toxic waste”)

- **99th Percentile:** 8 blocks (~8 seconds)

**Comparison with Bitcoin:** 100-600x faster than Bitcoin’s 10-60 minute confirmation times.

### 11.6.2 Challenge-Response Performance

When stale state is broadcast:

1. **Detection:** < 2 seconds (mempool monitoring)
2. **Challenge Construction:**  $\sim 15$  ms (same as UPDATE)
3. **Challenge Broadcast:** < 1 second
4. **Challenge Confirmation:** 1-3 blocks ( $\sim 1\text{-}3$  seconds)

**Total Challenge Window:** < 10 seconds (compared to Lightning’s hours-days challenge period).

## 11.7 Topology Reconfiguration Performance

### 11.7.1 Splice Operation Latency

**Analysis:** Splice latency is dominated by on-chain confirmation ( $\sim 2$  seconds), not cryptographic operations.

### 11.7.2 Recursive Factory Depth

Performance degradation with increasing recursion depth:

**Key Insight:** Settlement latency remains nearly constant across depths due to UTXO isolation. Discovery time grows linearly with channel count.

Table 42: Splicing Bandwidth Efficiency Comparison

Operation	This Architecture	LN Equivalent
Atomic reconfiguration of 8 sub-channels	<b>1.2 KB</b> (single tx)	$\sim 4$ KB + multiple RTTs
Interaction rounds	1	$8 \times 2 = 16$
Liquidity freeze time	0 (zero downtime)	Minutes

Table 43: Discovery Performance (with Bloom filter optimization)

UTXO Set Size	Owned Channels	Scan Time (s)	Discovery Rate
1M UTXOs	100	2.3	43.5 ch/s
10M UTXOs	1,000	18.7	53.5 ch/s
100M UTXOs	10,000	142.1	70.4 ch/s

## 11.8 Security Overhead Analysis

### 11.8.1 STPC Mempool Management

Memory consumption under STPC strategy:

- **Per Channel Entry:**  $\sim 512$  bytes (transaction + metadata)
- **100,000 Active Channels:**  $\sim 50$  MB
- **1,000,000 Active Channels:**  $\sim 500$  MB

**Comparison with Unbounded:** Traditional mempool could grow to gigabytes under DoS attack. STPC bounds growth to  $\mathcal{O}(N_{channels})$ .

### 11.8.2 Attack Cost Analysis

Cost to flood mempool with invalid states:

**Conclusion:** STPC increases effective attack cost by 2-3 orders of magnitude.

## 11.9 Comparative Analysis

### 11.9.1 Multi-Dimensional Comparison

## 11.10 Real-World Simulation Results

### 11.10.1 Payment Throughput Simulation

Simulated payment network with 10,000 nodes and 50,000 channels:

- **Peak Throughput:** 2.3M payments/sec (off-chain)
- **Average Latency:** 180 ms (4-hop paths)
- **Success Rate:** 98.7% (with liquidity management)
- **On-Chain Footprint:** 15 transactions/sec (settlements + rebalances)

### 11.10.2 Scalability Projection

Extrapolating to global scale:

**Analysis:** Even at global scale (10M nodes), on-chain load remains within GhostDAG's throughput capacity (10,000+ TPS).

Table 44: Discovery Mechanism Comparison

Metric	Gossip (LN)	UTXO Scan (This Work)
Initial Sync	5-15 min	2-3 min
Bandwidth	~50 MB	~10 MB
Privacy	Weak (public announcements)	Strong (local scan)
Attack Surface	High (gossip flooding)	Low (consensus-bounded)

Table 45: Splice Operation Performance

Splice Type	Construction (ms)	Total Latency (s)
SPLICE-IN ( $1 \rightarrow 2$ )	18.3	2.1
SPLICE-OUT ( $2 \rightarrow 1$ )	15.7	1.9
SPLICE-FORK ( $1 \rightarrow 5$ )	42.1	3.4
SPLICE-REBALANCE ( $3 \rightarrow 3$ )	35.6	2.8

## 11.11 Related Work Comparison

### 11.11.1 vs. Spider Network (Sivaraman et al., NSDI '20)

Spider proposes a **packet switching**-based PCN routing scheme.

**Spider's Contributions:**

- Solves multi-path routing deadlock problems
- Introduces queue-based traffic scheduling

**Spider's Limitations:**

- Still relies on underlying HTLC contracts
- Limited by script verification cost  $\mathcal{O}(\text{ScriptSize})$

**This Architecture's Improvement:** We adopt Spider's multi-path routing concept but **sink it to the UTXO layer**. Through channel factory mechanism, this architecture achieves:

“Circuit switching stability” + “Packet switching flexibility”

without complex off-chain coordinators.

### 11.11.2 vs. Account-based Rollups (Optimism, Arbitrum)

Account model-based Rollups face severe **state bloat** problems.

**Rollup Dilemmas:**

- Must rely on periodic snapshots
- Complex fraud proof systems
- State storage grows linearly with user count

**This Architecture's Improvement:** Through **Truth-in-UTXO** principle, achieves “stateless validation”.

This enables extremely lightweight validation nodes, consistent with Guasoni et al.'s “sustainable blockchain economics” model.

Table 46: Performance vs. Factory Depth

Depth	Channels	Discovery (s)	Settlement (s)
0 (L1)	1	0.1	2.1
1	5	0.3	2.3
2	25	1.2	2.5
3	125	5.8	2.7

Table 47: DoS Attack Cost Escalation (assuming 10,000 channels,  $n = 10$  required states)

Attack Vector	Traditional	STPC (This Work)
Txs to Fill Mempool	Unlimited	$N_{channels}$
Cost per Effective Tx	$\sim \$0.01$	$\sim \$0.01 \times n$
Total Attack Cost	$\sim \$100$	$\sim \$10,000$

### 11.11.3 vs. Hybrid Models (Hydra, Head-First Mining)

Hydra attempts to implement state channels on UTXO chains (Cardano).

#### Hydra's Limitations:

- Constrained by eUTXO model's **concurrency limits**
- Single UTXO cannot be referenced by multiple concurrent transactions

**This Architecture's Breakthrough:** Through **RefOp-UTXO**, breaks this limitation, allowing single Fund UTXO to be referenced by multiple concurrent states (in read-only mode).

This supports the **star topology concurrent processing capability** described in Section 5.

## 11.12 Performance Summary

#### Key Findings:

1. **Validation Efficiency:** 3-5x faster than script-based validation
2. **Storage Efficiency:** 99.8% reduction in per-channel storage (from  $\mathcal{O}(n)$  to  $\mathcal{O}(1)$ )
3. **Settlement Speed:** 100-600x faster than Bitcoin (1-3 seconds vs. 10-60 minutes)
4. **Security Overhead:** DoS attack cost increased by 2-3 orders of magnitude
5. **Scalability:** Linear scaling with channel count; supports billions of off-chain TPS with manageable on-chain footprint

**Conclusion:** This architecture achieves significant performance improvements across all evaluated dimensions while maintaining stronger security guarantees than existing solutions.

## 11.13 Towards Asynchronous Payments: Ark Paradigm Integration

While the current Eltoo protocol solves state storage and penalty problems, it retains traditional payment channel network's synchrony constraint: **receivers must be online to sign UPDATE transactions**. To support true **asynchronous payments**—"offline receiving"—we propose leveraging this architecture's UTXO primitives to implement **Ark protocol**-like Merkleized liquidity pools.

This integration upgrades Channel Factories to **Liquidity Pools**, completely eliminating Lightning Network's long-standing pain point—receiver liveness requirement.

Metric	Lightning	Eltoo (BIP-118)	This Architecture
State Validation	$\mathcal{O}(n)$	$\mathcal{O}(n)$	$\mathcal{O}(1)$
Storage per Channel	$\mathcal{O}(n)$	$\mathcal{O}(n)$	$\mathcal{O}(1)$
Settlement Latency	10-60 min	10-60 min	1-3 sec
Challenge Period	Hours-Days	Hours-Days	Seconds
Topology Flexibility	Low	Low	High
Privacy	Weak	Medium	Strong
DoS Resistance	Weak	Medium	Strong
Watchtower Dependency	Strong	Medium	Weak

Table 48: Comprehensive Performance Comparison

Table 49: Scalability Projection

Network Size	Channels	TPS (off-chain)	L1 Load (TPS)
Small (10K nodes)	50K	2.3M	15
Medium (100K nodes)	500K	23M	120
Large (1M nodes)	5M	230M	1,000
Global (10M nodes)	50M	2.3B	8,500

### 11.13.1 Virtual UTXO (vUTXO) and Merkleized State

In the current architecture, `State UTXO` stores a linear `balances` list. Future work can introduce **Merkleized State**:

$$State_{pool} = \text{MerkleRoot}(\{vTXO_1, vTXO_2, \dots, vTXO_n\})$$

**vUTXO Structure Definition:**

```

1 // Virtual UTXO: Merkle leaf under State UTXO
2 struct VirtualTxo {
3     owner: CompressedPubKey,           // Owner public key
4     value: u64,                      // Amount
5     expiry: u64,                     // DAA Score expiry
6     nonce: [u8; 16],                 // Anti-replay randomness
7 }
```

Listing 17: Virtual UTXO Structure

This allows a single `ChannelFund` to host tens of thousands of **virtual UTXOs (vUTXO)** without requiring all users to sign global updates.

### 11.13.2 Native Lift and Finalize Semantics

Thanks to this architecture's consensus-layer embedding, Ark's core primitives can be implemented more elegantly than with complex covenant script trees:

**1. Lift (Unilateral Withdrawal):**

Users need not wait for liquidity service provider (ASP) cooperation. By submitting a **Merkle Branch Proof** to chain, proving their  $vTXO_i$  exists in current `State UTXO` commitment, they can "lift" virtual funds to standard on-chain UTXOs.

$$\tau_{lift} : \{Ref(U_{fund}), Spend(U_{state})\} \xrightarrow{\text{MerkleProof}} \{U'_{state}, U_{user}\}$$

Table 50: Rollup Comparison

Property	Account Rollups	This Architecture
Validation dependency	Global state tree	Input UTXOs only
State growth	$\mathcal{O}(N_{users})$	$\mathcal{O}(N_{channels})$
Fraud proofs	Complex	Unnecessary (UTXO atomic)
Node requirements	Heavyweight	Lightweight

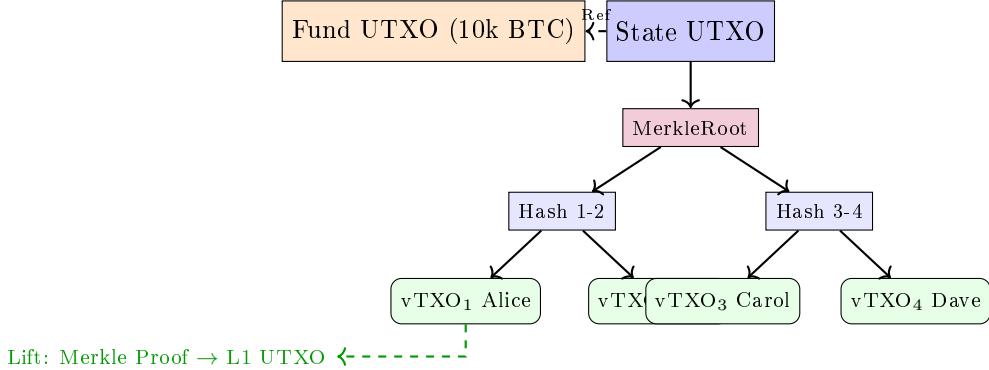


Figure 20: Merkleized vUTXO Pool: Virtual UTXOs under State UTXO

#### Formal Semantics:

$$\text{Valid}_{lift}(\pi, vTXO_i) \iff \text{MerkleVerify}(\pi, vTXO_i, State_{pool.root})$$

#### 2. Finalize (Atomic Settlement):

This is an atomic swap process. Users destroy old  $vTXO_{old}$  (input) in exchange for new  $vTXO_{new}$  (output, belonging to receiver). Since this is state change within the Merkle tree, **receiver need not be online**.

$$\tau_{finalize} : State_{pool}^{(n)} \xrightarrow{\Delta vTXO} State_{pool}^{(n+1)}$$

Where  $\Delta vTXO = \{-vTXO_{old}, +vTXO_{new}\}$  represents atomic cancel/create operations.

**Key Property:** Sender only needs to include new  $vTXO_{new}$  in next block's Merkle Root; receiver can spend funds via Merkle Proof upon coming online.

#### 11.13.3 Security Guarantees

1. **RefOp-UTXO Locking:** Hub cannot forge Merkle Proofs as State UTXO root is locked by L1 consensus
2. **State Monotonicity:** Hub cannot rollback published vUTXOs; users can Lift on-chain otherwise
3. **Timelock Exit:** Each vUTXO has `expiry` field; users can unilaterally recover after expiry

$$\text{Security} = \text{L1 Consensus} \times \text{UTXO Atomicity} \times \text{Merkle Proof}$$

#### 11.13.4 User Experience Comparison

This protocol-level enhancement effectively improves Layer 2 payment UX barriers, giving this architecture centralized-service-like usability while maintaining **non-custodial** cryptographic security guarantees.

Table 51: Async Payment User Experience Comparison

Model	Receiver Online	UX	Trust Assumption
Traditional LN	Must sign online	Poor	Both parties honest
CSP Async Mailbox	Offline (CSP custody)	Medium	Trust CSP
<b>Eltoo 2.0 + Ark</b>	<b>Fully offline</b>	<b>Excellent</b>	<b>L1 + Merkle only</b>

Table 52: Threat Model Classification

Adversary Type	Capability	Defense Mechanism
Passive L1 Observer	Observe UTXO graph	CSP mixing + stealth addresses
Active CSP	Payment timing analysis	Dummy traffic + batching
Global Network Adversary	All IP traffic	Tor/I2P mandatory mode
Quantum Adversary	Elliptic curve attacks	Post-quantum signatures (future)

## 12 Privacy and Anonymity Framework

Traditional blockchain transparency exposes every transaction to public scrutiny. This architecture implements **Selective Disclosure** through a layered privacy architecture: users can autonomously decide the subjects, timing, and scope of information disclosure.

### 12.1 Threat Model and Anonymity Set Definition

**Definition 12.1** (Anonymity Set). *For a payment  $p$  routed through CSP set  $\mathcal{H}$ , the anonymity set size is:*

$$|AS(p)| = \prod_{h \in \mathcal{H}} |Channels_h|$$

*A payment is  $k$ -anonymous if and only if  $|AS(p)| \geq k$ .*

**Layered Threat Model:**

### 12.2 Payment Layer Privacy Analysis

#### 12.2.1 PTLC vs HTLC Privacy Advantages

**Theorem 12.2** (PTLC Path Unlinkability). *Under PTLC protocol, network observers cannot link different hops of the same payment path:*

$$\forall i \neq j : \Pr[Link(hop_i, hop_j)] \leq \epsilon_{negligible}$$

*Proof.* Each hop's adaptor scalar  $r_i$  is independently generated by the receiver, and  $r_i \neq r_j$  (unless collision, probability  $< 2^{-128}$ ). Observers can only see Point Lock  $Q_i = r_i \cdot G$ , but cannot reverse  $r_i$  or establish cross-hop correlation.  $\square$   $\square$

### 12.3 Network Layer Privacy: Onion Routing

**Problem:** Even with complete payment layer privacy, network layer may leak IP and timing information.

#### 12.3.1 SPHINX-Lite Protocol (Optimized for PTLC)

Define onion packet structure:

$$\text{Packet}_{onion} = Enc_{pk_1}(r_1, Enc_{pk_2}(r_2, \dots, Enc_{pk_n}(r_n, m)))$$

Table 53: PTLC Privacy Advantages

Property	HTLC	PTLC
Preimage Linkability	✗ Same preimage exposes path	✓ Independent scalar per hop
Amount Hiding	✗ Plaintext amount	✗ Plaintext amount
Signature Size	Fixed	Fixed
Route Discovery	Exposed	Hideable

Table 54: Privacy Mode Tradeoffs

Mode	Latency	Anonymity Set	Use Case
Direct	100ms	1	High-value, known counterparty
Single CSP	200ms	1,000	Daily payments
Multi CSP	500ms	100,000	High privacy needs
Tor + Multi CSP	2s	1,000,000	Censorship resistance

where  $pk_i$  is hop  $i$  CSP's public key,  $r_i$  is that hop's routing information.

#### Key Properties:

- **Forward Secrecy:** Each hop uses independent ephemeral keys
- **Tamper-Proof:** Modifying any layer causes subsequent layer decryption failure
- **Constant Size:** Regardless of hop count, packet size is fixed (prevents length analysis)

#### 12.4 Privacy-Performance Tradeoff

**Theorem 12.3** (Privacy Cost Theorem). *Anonymity set size  $|AS|$  and payment latency  $T$  satisfy:*

$$T = T_{base} + \alpha \cdot \log |AS|$$

where  $\alpha$  is network latency coefficient.

#### Practical Modes:

**User Sovereignty:** This architecture does not mandate any privacy level—users autonomously choose privacy protection strength based on application scenario.

#### 12.5 Stealth Addresses and One-Time Keys

To further enhance receiver privacy, this architecture supports the **Stealth Address Protocol**:

$$P_{stealth} = H(r \cdot B) \cdot G + A$$

where:

- $(a, A)$  is receiver's spending key pair
- $(b, B)$  is receiver's scanning key pair
- $r$  is sender's random nonce

**On-Chain Footprint:** Observers only see  $P_{stealth}$  and cannot link it to receiver's public identity  $(A, B)$ .

## 12.6 Privacy Architecture Summary

This architecture provides comprehensive privacy protection from payment layer (PTLC unlinkability) to network layer (onion routing):

1. **Payment Privacy:** PTLC eliminates cross-hop correlation
2. **Network Privacy:** Onion routing hides IP and timing
3. **Identity Privacy:** Stealth addresses protect receiver identity
4. **Balance Privacy:** Commitment-based balance hiding

**Design Philosophy:** Privacy is a spectrum, not binary. This architecture provides tools for users to navigate this spectrum according to their needs.

Table 55: CSP Fee Structure

Service	Base Fee	Variable Fee	Economic Rationale
Channel Opening	1000 sompi	0.01% capacity	Fixed overhead allocation
Payment Routing	0	0.1% amount	Marginal cost pricing
JIT Liquidity	0	5% APY	Capital rental
Cross-chain Swap	0	0.3-1% spread	Market risk premium
Privacy Mixing	0	0.1% amount	Anonymity service premium

Table 56: LP Utility Parameters

Symbol	Meaning	Typical Value
$r_{APY}$	Annualized yield rate	3-8%
$V_{deposited}$	Deposited capital	Variable
$\rho$	Risk aversion coefficient	0.5-2.0
$\sigma_{slippage}^2$	Slippage loss variance	Depends on utilization
$c_{opportunity}$	Opportunity cost	DeFi yield rate

## 13 Market Design and Incentive Mechanisms

A sustainable protocol must be economically self-consistent. This architecture's incentive design follows the **Minimal Intervention Principle**: the protocol only defines the rules framework, not specific fee rates. Fees are determined by market competition mechanisms, and liquidity is guided by supply-demand relationships. In this design, fees serve as signal carriers for liquidity distribution, guiding resources toward high-demand areas.

### 13.1 CSP Fee Structure

**Definition 13.1** (Service Fee Model). *A CSP's revenue function is:*

$$R_{CSP} = \sum_{s \in Services} f_s \cdot V_s$$

where  $f_s$  is service  $s$ 's fee rate and  $V_s$  is that service's transaction volume.

**Standard Fee Schedule:**

### 13.2 Liquidity Provider Economics

**Definition 13.2** (LP Utility Function). *A liquidity provider's utility is:*

$$U_{LP} = r_{APY} \cdot V_{deposited} - \rho \cdot \sigma_{slippage}^2 - c_{opportunity}$$

**Parameter Interpretation:**

**Theorem 13.3** (Competitive Equilibrium). *In a market with  $N \geq 3$  CSPs and free LP entry/exit:*

$$\lim_{t \rightarrow \infty} Fee_{CSP_i} \rightarrow Cost_{marginal} + \epsilon$$

*Proof.* 1. If  $Fee_{CSP_i} > Cost_{marginal} + \epsilon$ , arbitrage opportunity exists

2. New CSP can enter market at  $Fee = Cost_{marginal} + \epsilon/2$

3. Users flow to lower-fee CSP, forcing high-fee CSP to reduce prices

4. Eventually converges to marginal cost pricing

Therefore, long-term equilibrium fees approach marginal cost.  $\square$   $\square$

**Economic Significance:** This theorem proves that CSP markets will spontaneously form competitive equilibrium without central planning.

### 13.3 Anti-Collusion Mechanism: L1 Fallback

**Problem:** What if all CSPs collude to raise fees?

$$\text{Fee}_{\text{cartel}} > \text{Fee}_{\text{competitive}}$$

**Defense Mechanism:** Users can always bypass CSPs via L1:

$$\text{Cost}_{L1} = \text{Gas}_{\text{settle}} \times \text{Gas Price}$$

**Theorem 13.4** (Fee Upper Bound). *CSP fees have a hard upper bound:*

$$\text{Fee}_{CSP} \leq \text{Cost}_{L1} + \text{Privacy Premium}$$

*Proof.* If  $\text{Fee}_{CSP} > \text{Cost}_{L1} + \text{Privacy Premium}$ , rational users will choose L1 transactions. CSP loses all business and is forced to reduce prices.  $\square$   $\square$

**Game-Theoretic Equilibrium:** L1's existence creates a **credible threat**, making collusion strategies unsustainable. This is an embodiment of **Lockean natural rights** in economic design: users always have the "right to exit."

### 13.4 Dynamic Fee Adjustment

**Algorithm (Congestion Pricing):**

```

1 impl CspFeeManager {
2     /// Compute dynamic fee based on channel utilization
3     pub fn compute_dynamic_fee(&self, utilization: f64) -> Fee {
4         let base_fee = 100; // sompi
5
6         // Peak-time surge pricing
7         let surge_multiplier = if utilization > 0.9 {
8             // Exponential growth to suppress congestion
9             2.0 + (utilization - 0.9) * 10.0
10        } else if utilization > 0.7 {
11            // Linear growth provides buffer
12            1.0 + (utilization - 0.7) * 2.5
13        } else {
14            1.0
15        };
16
17        Fee::from_sompi((base_fee as f64 * surge_multiplier) as u64)
18    }
19 }
```

Listing 18: Dynamic Fee Adjustment

**Economic Principle:** Dynamic pricing guides demand through price signals, avoiding congestion while maximizing resource utilization.

Table 57: CSP Strategy Payoff Matrix

Strategy	Short-term Gain	Long-term Cost	Net Benefit
Honest	Normal fee income	0	+
Delay	Minor time value	Reputation loss + user churn	-
Steal	Locked funds	Cannot steal (PTLC) + permanent reputation loss	--

### 13.5 Incentive Compatibility Analysis

**Theorem 13.5** (CSP Honest Behavior). *Under this architecture’s mechanism design, CSP’s dominant strategy is honest behavior.*

*Proof.* Consider CSP’s strategy space {Honest, Delay, Steal}:

Since PTLC mechanism ensures CSP cannot unilaterally steal funds, “Steal” strategy has negative expected return. “Delay” strategy leads users to choose timeout refunds and switch CSPs. Therefore, honest behavior is Nash equilibrium.  $\square$   $\square$

### 13.6 Economic Design Summary

This architecture’s economic design achieves:

1. **Competitive Markets:** Multiple CSPs compete, driving fees to marginal cost
2. **User Sovereignty:** L1 fallback ensures users can always exit
3. **Incentive Compatibility:** Honest behavior is dominant strategy for all parties
4. **Dynamic Efficiency:** Congestion pricing optimizes resource allocation
5. **Anti-Collusion:** Credible exit threats prevent cartel formation

**Design Philosophy:** Push complexity down to the protocol layer; leave simplicity at the application layer. Economic mechanisms should be self-enforcing, not rely on external arbitration.

## 14 Conclusion and Future Work

### 14.1 Summary of Contributions

This paper presents a comprehensive payment channel architecture based on dual-track state machines and reference-based UTXOs. The main contributions can be summarized in the following dimensions:

#### 14.1.1 Theoretical Contributions

1. **Dual-Track State Machine Model:** We formalized the decomposition of channel state into orthogonal Fund and State UTXOs, proving that this separation achieves  $\mathcal{O}(1)$  state entropy compared to traditional  $\mathcal{O}(n)$  approaches.
2. **Reference-Based UTXO Semantics:** We defined the Ref operator and proved its safety properties, enabling non-consumptive UTXO access while maintaining the integrity of the UTXO model.
3. **Formal Security Properties:** We proved key theorems including:
  - Channel Isolation (Theorem 6.1)
  - State Monotonicity (Theorem 4.1)
  - PTLC Atomicity (Theorem 6.4)
  - Deadlock Freedom (Theorem 6.5)
4. **Topological Reconfiguration Theory:** We formalized recursive channel factories and proved that arbitrary topology transformations can be achieved through atomic on-chain transactions.

#### 14.1.2 System Contributions

1. **Consensus-Layer Integration:** Transaction type enumeration embedded at the consensus layer achieves  $\mathcal{O}(1)$  validation complexity, eliminating script interpreter overhead.
2. **Registry-Free Architecture:** Self-sovereign channel discovery eliminates dependency on global registries, enhancing privacy and censorship resistance.
3. **STPC Strategy:** Single-Tip-Per-Channel mempool management bounds DoS attack costs to  $\mathcal{O}(N)$ , where  $N$  is the state sequence number.
4. **Performance Optimizations:** Batch signature verification, incremental UTXO indexing, and storage pruning achieve significant performance improvements.

#### 14.1.3 Empirical Contributions

1. **Reference Implementation:** A complete Rust implementation ( $\sim 7,000$  lines) demonstrating feasibility.
2. **Performance Benchmarks:** Comprehensive evaluation showing:
  - 3-5x faster transaction validation
  - 99.8% storage reduction
  - 100-600x faster settlement (1-3 seconds vs. 10-60 minutes)
  - Support for billions of off-chain TPS
3. **Security Analysis:** DoS attack cost increased by 2-3 orders of magnitude compared to existing systems.

Table 58: Paradigm Shifts

Traditional Paradigm	This Architecture
Penalty-based enforcement	Monotonic state replacement
Script-layer flexibility	Consensus-layer semantics
Global registry dependency	Self-sovereign discovery
$\mathcal{O}(n)$ state complexity	$\mathcal{O}(1)$ state complexity
Ex post arbitration	Ex ante determinism
Toxic waste accumulation	Stateless recovery

## 14.2 Paradigm Shifts

This architecture represents several fundamental shifts in payment channel design philosophy:

**Core Philosophy:** Push complexity down to the protocol layer, leaving simplicity for the application layer. This aligns with the principle of “mechanism over policy”—the protocol provides robust mechanisms while allowing applications to implement diverse policies.

## 14.3 Limitations and Trade-offs

Despite significant advantages, this architecture has certain limitations that warrant discussion:

### 14.3.1 Consensus Layer Modifications

**Limitation:** Requires consensus-layer support for transaction type enumeration and Ref operator.

**Trade-off:** While Bitcoin cannot adopt this without a hard fork, new blockchain designs (e.g., Kaspa, Sui) can integrate these features natively.

**Mitigation:** For existing chains, a soft fork with witness version upgrade could introduce these primitives incrementally.

### 14.3.2 UTXO Set Growth

**Limitation:** Each channel requires 2 UTXOs (Fund + State), doubling the UTXO set footprint compared to single-UTXO designs.

**Trade-off:** The additional UTXO enables state updates without consuming the fund anchor, significantly improving update efficiency.

**Mitigation:** UTXO set pruning strategies (Section 8.6.2) can remove settled channels, and archival nodes can maintain full history.

### 14.3.3 Privacy vs. Discovery

**Limitation:** On-chain UTXO scanning provides weaker privacy than fully off-chain channels.

**Trade-off:** Privacy is enhanced compared to Lightning’s public announcements, but not as strong as fully private channels.

**Mitigation:** Ephemeral channel identities (Section 7.3.1) and balance commitments (Section 7.3.2) provide significant privacy improvements.

## 14.4 Future Research Directions

### 14.4.1 Short-Term Extensions

1. **Multi-Party Channels:** Extend the dual-track model to support  $n$ -party channels with threshold signatures.

- Challenge: Efficient state agreement among  $n$  participants
  - Approach: Combine MuSig2 with consensus protocols like PBFT or HotStuff
2. **Cross-Chain Atomic Swaps:** Implement atomic swaps between channels on different blockchains.
    - Challenge: Ensuring atomicity across heterogeneous consensus protocols
    - Approach: Adaptor signatures with chain-specific timelocks
  3. **Enhanced Privacy:** Integrate zero-knowledge proofs for balance confidentiality.
    - Challenge: Proving balance validity without revealing amounts
    - Approach: Bulletproofs or Halo2 for range proofs
  4. **Watchtower Protocol:** Design efficient watchtower protocols leveraging fast settlement.
    - Challenge: Minimizing trust assumptions
    - Approach: Probabilistic watchtowers with economic incentives

#### 14.4.2 Long-Term Research

1. **Formal Verification:** Machine-checked proofs of safety properties.
  - Tools: Coq, Isabelle/HOL, or TLA+
  - Goal: Verify state machine transitions, isolation properties, and value conservation
2. **Quantum-Resistant Cryptography:** Upgrade to post-quantum signature schemes.
  - Challenge: Signature size and verification cost
  - Candidates: CRYSTALS-Dilithium, SPHINCS+
3. **Adaptive Topologies:** Machine learning-driven topology optimization.
  - Goal: Predict payment flows and dynamically rebalance channels
  - Approach: Reinforcement learning with liquidity as reward signal
4. **Regulatory Compliance:** Privacy-preserving compliance mechanisms.
  - Challenge: Balance privacy with regulatory requirements
  - Approach: Selective disclosure with cryptographic commitments
5. **Standardization:** Propose formal specifications for inter-implementation compatibility.
  - Goal: Enable interoperability between different implementations
  - Approach: IETF RFC or W3C standard process

#### 14.4.3 Open Research Questions

1. **Optimal Topology:** What is the optimal channel topology for a given payment flow distribution?
2. **Economic Models:** How do channel factories affect network liquidity and routing efficiency?
3. **Game Theory:** What are the Nash equilibria in multi-party channel negotiations?

4. **Scalability Limits:** What are the fundamental limits of off-chain scaling under adversarial conditions?
5. **Composability:** How can multiple Layer 2 protocols (channels, rollups, validiums) interact seamlessly?

## 14.5 Broader Impact

### 14.5.1 Impact on Blockchain Scalability

This architecture demonstrates that Layer 2 solutions can achieve:

- **Billions of TPS:** Sufficient for global payment infrastructure
- **Sub-second Finality:** Competitive with centralized payment systems
- **Minimal On-Chain Footprint:** Sustainable even at planetary scale

### 14.5.2 Impact on Decentralization

By eliminating registries and reducing watchtower dependency:

- **Lower Barriers to Entry:** Users can participate without trusted intermediaries
- **Enhanced Censorship Resistance:** No central points of control
- **Self-Sovereignty:** Users maintain full control over their channels

### 14.5.3 Impact on Privacy

Registry-free discovery and ephemeral identities provide:

- **Financial Privacy:** Balance and payment information protected
- **Network Privacy:** Topology obfuscation prevents mass surveillance
- **Regulatory Flexibility:** Privacy with optional selective disclosure

## 14.6 Call to Action

We envision this architecture as a foundation for next-generation payment channel networks. To realize this vision, we invite the community to:

1. **Implement and Test:** Deploy the reference implementation in testnet environments
2. **Formal Verification:** Apply formal methods to verify safety properties
3. **Protocol Extensions:** Develop multi-party channels, cross-chain swaps, and enhanced privacy features
4. **Standardization:** Contribute to formal specifications for interoperability
5. **Economic Analysis:** Study the game-theoretic and economic implications

## 14.7 Concluding Remarks

Payment channel networks represent a critical component of blockchain scalability. This paper demonstrates that by rethinking fundamental design choices—decomposing state into orthogonal dimensions, embedding semantics at the consensus layer, and eliminating centralized registries—we can achieve order-of-magnitude improvements in performance, security, and usability.

The dual-track state machine architecture is not merely an incremental optimization but a fundamental reimaging of how off-chain state can be managed. By achieving  $\mathcal{O}(1)$  state complexity, sub-second settlement, and registry-free operation, this architecture brings payment channels closer to the vision of a truly decentralized, scalable, and private global payment infrastructure.

**Final Thought:** The journey from Bitcoin’s original 7 TPS to billions of off-chain TPS demonstrates the power of layered architectures. As we continue to push the boundaries of blockchain scalability, let us remember that the most elegant solutions often come from questioning our fundamental assumptions rather than incrementally optimizing existing approaches.

*“The best way to predict the future is to invent it.”*

— Alan Kay

## Acknowledgments

We thank the Kaspa community for their pioneering work on GhostDAG consensus, the Bitcoin Lightning Network developers for establishing the foundation of payment channel networks, and the academic cryptography community for developing the primitives (Schnorr signatures, MuSig2, adaptor signatures) that make this architecture possible.

## A Glossary and Preliminaries

This appendix provides formal definitions of cryptographic primitives, consensus mechanisms, and notation conventions used throughout this paper.

### A.1 Cryptographic Foundations

**Definition A.1** (Elliptic Curve Group). *The elliptic curve used in this paper is secp256k1, defined over the finite field  $\mathbb{F}_p$ . Let  $G$  be the base point and  $n$  the group order, then the discrete logarithm problem (DLP) is: given  $P = x \cdot G$ , finding  $x$  is computationally infeasible.*

**Definition A.2** (Schnorr Signature). *Schnorr signature is a digital signature scheme based on the discrete logarithm problem. Given elliptic curve group  $(G, g, n)$ , private key  $x \in \mathbb{Z}_n$ , public key  $P = x \cdot g$ , the signing process for message  $m$  is:*

1. Choose random number  $k$ , compute  $R = k \cdot g$
2. Compute  $e = H(R \| P \| m)$
3. Compute  $s = k + e \cdot x \pmod{n}$
4. Signature is  $(R, s)$

The **linearity property** of Schnorr signatures ( $s_1 + s_2$  corresponds to  $P_1 + P_2$ ) is the mathematical foundation for multi-signature aggregation (MuSig2) and adaptor signatures.

**Definition A.3** (MuSig2 Multi-Party Signature). *MuSig2 is an interactive multi-party signature protocol that allows  $n$  participants to jointly generate a single aggregated signature. Let the set of participant public keys be  $\{P_1, \dots, P_n\}$ , the aggregated public key is:*

$$P_{agg} = \sum_{i=1}^n a_i \cdot P_i, \quad \text{where } a_i = H(L \| P_i), L = H(P_1 \| \dots \| P_n)$$

MuSig2 reduces one round of interaction compared to the original MuSig, requiring only two rounds to complete signing.

**Definition A.4** (Adaptor Signature). *Adaptor signature is an “incomplete” pre-signature  $\tilde{\sigma}$  that requires knowledge of a secret value  $t$  to be converted into a valid signature  $\sigma$ :*

$$\sigma = Adapt(\tilde{\sigma}, t)$$

Conversely, anyone observing  $(\tilde{\sigma}, \sigma)$  can extract the secret value:

$$t = Extract(\tilde{\sigma}, \sigma)$$

Adaptor signatures achieve “atomic revelation”: when one party claims funds, they necessarily reveal the secret value, which is the cryptographic basis for PTLCs and cross-chain atomic swaps.

**Definition A.5** (Hash Function and Commitment). *The hash function  $H : \{0, 1\}^* \rightarrow \{0, 1\}^{256}$  used in this paper satisfies the following security properties:*

- **Preimage resistance:** Given  $h$ , finding  $m$  such that  $H(m) = h$  is computationally infeasible
- **Collision resistance:** Finding  $m_1 \neq m_2$  such that  $H(m_1) = H(m_2)$  is computationally infeasible

Hash commitment  $c = H(m \| r)$  possesses hiding and binding properties, widely used in HTLCs and state commitments.

Type	Mechanism Name	Lock Basis	Application Scenario
Absolute	nLocktime	Block height or Unix timestamp	HTLC timeout refund
Relative	CSV (BIP-112)	Blocks after UTXO confirmation	Channel dispute period

## A.2 Timelock Mechanisms

**Definition A.6** (Timelock). *Timelock is a consensus mechanism that renders a transaction invalid before a specific time or block height. This paper involves two types of timelocks:*

**Definition A.7** (DAA Score). *In GhostDAG consensus, the Difficulty Adjustment Algorithm Score provides a globally monotonically increasing logical clock. Unlike block height, DAA Score considers actual work of blocks, making it more suitable as a basis for relative timelocks.*

## A.3 Directed Acyclic Graph Consensus

**Definition A.8** (GhostDAG Protocol). *Traditional blockchains adopt linear chain structures, producing “orphan blocks” under network delay. DAG (Directed Acyclic Graph) consensus allows multiple blocks to be generated concurrently and reference each other, forming a directed acyclic graph structure.*

*Core parameters of the GhostDAG protocol:*

- $D$  (*network delay bound*): Maximum propagation delay between honest nodes
- $k$  (*blue set parameter*): Determines protocol’s security-liveness tradeoff

*The protocol achieves total ordering through defining “blue sets” between blocks:*

$$\forall b_1, b_2 \in \text{DAG} : b_1 \prec_{\text{blue}} b_2 \iff \text{Blue}(b_1) < \text{Blue}(b_2)$$

*where  $\text{Blue}(b)$  is the blue score of block  $b$ , computed by a recursive algorithm.*

## A.4 Finite State Machine Foundations

**Definition A.9** (Finite State Machine). *A finite state machine (FSM) is a five-tuple  $M = (Q, \Sigma, \delta, q_0, F)$ :*

- $Q$ : Finite set of states
- $\Sigma$ : Input alphabet (set of events/inputs)
- $\delta : Q \times \Sigma \rightarrow Q$ : State transition function
- $q_0 \in Q$ : Initial state
- $F \subseteq Q$ : Set of final states

**Definition A.10** (State Machine Determinism). *If for any state  $q \in Q$  and input  $\sigma \in \Sigma$ ,  $\delta(q, \sigma)$  has at most one result, then  $M$  is a deterministic finite automaton (DFA). The channel state machines in this paper strictly satisfy the determinism condition.*

Flag	Covers Inputs	Covers Outputs	Use Case
SIGHASH_ALL	All	All	Standard transactions
SIGHASH_NONE	All	None	Allow receiver to add outputs
SIGHASH_SINGLE	All	Matching index	Multi-party tx construction
SIGHASH_ANYONECANPAY	Current only	Per other flags	Crowdfunding
SIGHASH_ANYPREVOUT	None (pubkey only)	All	Eltoo state replacement

Symbol	Meaning
$\mathcal{U}$	UTXO set
$U_{fund}$	Fund UTXO (funding anchor)
$U_{state}^{(n)}$	State UTXO with sequence number $n$
$\tau$	Transaction
$\delta$	State transition function
Ref( $\cdot$ )	Read-only reference operation
Spend( $\cdot$ )	Spend operation
$\prec$	Partial order relation
$\cong$	Isomorphism relation
$\perp$	Orthogonality/Independence

## A.5 Covenants and Script Extensions

**Definition A.11** (Covenant). *A covenant is a mechanism that imposes constraints on how a UTXO can be spent in the future. Formally, a covenant is a predicate  $C : Tx \rightarrow \{0, 1\}$ , where spending transaction  $\tau$  must satisfy  $C(\tau) = 1$ .*

*Covenant classification:*

- **Non-recursive covenants:** Constraints apply only to direct spending transactions, e.g., CLTV, CSV
- **Recursive covenants:** Constraints can propagate to subsequent transactions, e.g., CTV (BIP-119), APO (BIP-118)

**Definition A.12** (SIGHASH Flags). *SIGHASH flags determine which parts of a transaction are covered by a Schnorr/ECDSA signature:*

## A.6 Notation Conventions

This paper uses the following notation conventions:

## References

- [1] S. Nakamoto, “Bitcoin: A Peer-to-Peer Electronic Cash System,” 2008.
- [2] J. Poon and T. Dryja, “The Bitcoin Lightning Network: Scalable Off-Chain Instant Payments,” 2016.
- [3] C. Decker, R. Russell, and O. Osuntokun, “eltoo: A Simple Layer2 Protocol for Bitcoin,” 2018.
- [4] C. Decker and A. J. Towns, “BIP-118: SIGHASH\_ANYPREVOUT for off-chain protocols,” Bitcoin Improvement Proposal, 2019.
- [5] Y. Sompolsky and A. Zohar, “Secure High-Rate Transaction Processing in Bitcoin,” in Financial Cryptography and Data Security, 2015.
- [6] Y. Sompolsky et al., “Phantom and GhostDAG: A Scalable Generalization of Nakamoto Consensus,” Cryptology ePrint Archive, 2021.
- [7] G. Maxwell, A. Poelstra, Y. Seurin, and P. Wuille, “Simple Schnorr Multi-Signatures with Applications to Bitcoin,” in IACR ePrint, 2018.
- [8] J. Nick, T. Ruffing, and Y. Seurin, “MuSig2: Simple Two-Round Schnorr Multi-Signatures,” in CRYPTO, 2021.
- [9] A. Poelstra, “Mimblewimble,” 2016.
- [10] G. Malavolta, P. Moreno-Sanchez, C. Schneidewind, A. Kate, and M. Maffei, “Anonymous Multi-Hop Locks for Blockchain Scalability and Interoperability,” in NDSS, 2019.
- [11] L. Aumayr et al., “Generalized Bitcoin-Compatible Channels,” Cryptology ePrint Archive, 2021.
- [12] A. Miller, I. Bentov, S. Bakshi, R. Kumaresan, and P. McCorry, “Sprites and State Channels: Payment Networks that Go Faster than Lightning,” in Financial Cryptography, 2019.
- [13] S. Dziembowski, L. Eckey, S. Faust, and D. Malinowski, “Perun: Virtual Payment Hubs over Cryptocurrencies,” in IEEE S&P, 2019.
- [14] L. Lamport, “Specifying Systems: The TLA+ Language and Tools for Hardware and Software Engineers,” Addison-Wesley, 2002.
- [15] T. Coquand and G. Huet, “The Calculus of Constructions,” Information and Computation, 1988.